

# Biostat 203B Homework 5

Due Mar 20 @ 11:59PM

AUTHOR

Sophia Luo, 106409469

```
library(tidyverse)
```

— Attaching core tidyverse packages — tidyverse 2.0.0 —

```
✓ dplyr      1.1.4    ✓ readr      2.1.5
✓ forcats    1.0.0    ✓ stringr    1.5.1
✓ ggplot2    3.5.1    ✓ tibble     3.2.1
✓ lubridate  1.9.4    ✓ tidyr      1.3.1
✓ purrr      1.0.4
```

— Conflicts — tidyverse\_conflicts() —

```
✗ dplyr::filter() masks stats::filter()
```

```
✗ dplyr::lag()     masks stats::lag()
```

! Use the conflicted package (<<http://conflicted.r-lib.org/>>) to force all conflicts to become errors

```
library(tidymodels)
```

— Attaching packages — tidymodels 1.3.0 —

```
✓ broom      1.0.7    ✓ rsample     1.2.1
✓ dials      1.4.0    ✓ tune        1.3.0
✓ infer      1.0.7    ✓ workflows   1.2.0
✓ modeldata  1.4.0    ✓ workflowsets 1.1.0
✓ parsnip    1.3.1    ✓ yardstick   1.3.2
✓ recipes    1.2.0
```

— Conflicts — tidymodels\_conflicts() —

```
✗ scales::discard() masks purrr::discard()
```

```
✗ dplyr::filter()   masks stats::filter()
```

```
✗ recipes::fixed()  masks stringr::fixed()
```

```
✗ dplyr::lag()       masks stats::lag()
```

```
✗ yardstick::spec() masks readr::spec()
```

```
✗ recipes::step()   masks stats::step()
```

```
library(gtsummary)
```

```
library(ggplot2)
```

```
library(recipes) # Data Preprocessing
```

```
library(glmnet) # Logistic Regression
```

Loading required package: Matrix

Attaching package: 'Matrix'

The following objects are masked from 'package:tidyr':

expand, pack, unpack

Loaded glmnet 4.1-8

```
library(caret) # Random Forest
```

Loading required package: lattice

Attaching package: 'caret'

The following objects are masked from 'package:yardstick':

precision, recall, sensitivity, specificity

The following object is masked from 'package:purrr':

lift

```
library(ranger) # Random Forest  
library(xgboost) # XGBoost
```

Attaching package: 'xgboost'

The following object is masked from 'package:dplyr':

slice

```
library(stacks) # model stacking  
library(broom) # extract model coefficients
```

## Predicting ICU duration

Using the ICU cohort `mimiciv_icu_cohort.rds` you built in Homework 4, develop at least three machine learning approaches (logistic regression with enet regularization, random forest, boosting, SVM, MLP, etc) plus a model stacking approach for predicting whether a patient's ICU stay will be longer than 2 days. You should use the `los_long` variable as the outcome. Your algorithms can use patient demographic information (gender, age at ICU `intime`, marital status, race), ICU admission information (first care unit), the last lab measurements before the ICU stay, and first vital measurements during ICU stay as features. You are welcome to use any feature engineering techniques you think are appropriate; but make sure to not use features that are not available at an ICU stay's `intime`. For instance, `last_careunit` cannot be used in your algorithms.

### 1. Data preprocessing and feature engineering. [🔗](#)

```
# Load Data  
cohort_data <- readRDS("../hw4/mimiciv_shiny/mimic_icu_cohort.rds")
```

```

# Select Features & Target
mimiciv_icu_cohort <- cohort_data |>
  select(
    subject_id, hadm_id, stay_id,

    los_long, # Target variable (ICU stay > 2 days)

    # Demographics
    gender, age_intime, marital_status, race,

    # ICU admission
    first_careunit,

    # Lab measurements (Last before ICU)
    bicarbonate, chloride, creatinine, glucose, hematocrit,
    potassium, sodium, wbc,

    # Vital measurements (First in ICU)
    heart_rate, non_invasive_blood_pressure_diastolic,
    non_invasive_blood_pressure_systolic, respiratory_rate,
    temperature_fahrenheit
  ) |>
# Convert categorical variables to factors
mutate(
  gender = as.factor(gender),
  marital_status = as.factor(marital_status),
  race = as.factor(tolower(race))
) |>
drop_na(los_long) |>
print(width = Inf)

```

# A tibble: 94,444 × 22

	subject_id	hadm_id	stay_id	los_long	gender	age_intime	marital_status	race
	<int>	<int>	<int>	<fct>	<fct>	<int>	<fct>	<fct>
1	10000032	29079034	39553978	FALSE	F	52	WIDOWED	white
2	10000690	25860671	37081114	TRUE	F	86	WIDOWED	white
3	10000980	26913865	39765666	FALSE	F	76	MARRIED	black
4	10001217	24597018	37067082	FALSE	F	55	MARRIED	white
5	10001217	27703517	34592300	FALSE	F	55	MARRIED	white
6	10001725	25563031	31205490	FALSE	F	46	MARRIED	white
7	10001843	26133978	39698942	FALSE	M	76	SINGLE	white
8	10001884	26184834	37510196	TRUE	F	77	MARRIED	black
9	10002013	23581541	39060235	FALSE	F	57	SINGLE	other
10	10002114	27793700	34672098	TRUE	M	56	<NA>	other
	first_careunit			bicarbonate chloride				
	<fct>			<dbl> <dbl>				
1	Medical Intensive Care Unit (MICU)			25 95				
2	Medical Intensive Care Unit (MICU)			26 100				
3	Medical Intensive Care Unit (MICU)			21 109				

4	Surgical Intensive Care Unit (SICU)	22	108
5	Surgical Intensive Care Unit (SICU)	30	104
6	Medical/Surgical Intensive Care Unit (MICU/SICU)	NA	98
7	Medical/Surgical Intensive Care Unit (MICU/SICU)	28	97
8	Medical Intensive Care Unit (MICU)	30	88
9	Cardiac Vascular Intensive Care Unit (CVICU)	24	102
10	Other	18	NA

	creatinine	glucose	hematocrit	potassium	sodium	wbc	heart_rate
	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>
1	0.7	102	41.1	6.7	126	6.9	91
2	1	85	36.1	4.8	137	7.1	78
3	2.3	89	27.3	3.9	144	5.3	76
4	0.6	112	38.1	4.2	142	15.7	86
5	0.5	87	37.4	4.1	142	5.4	79.3
6	NA	NA	NA	4.1	139	NA	86
7	1.3	131	31.4	3.9	138	10.4	124.
8	1.1	141	39.7	4.5	130	12.2	49
9	0.9	288	34.9	3.5	137	7.2	80
10	3.1	95	34.3	6.5	125	16.8	110.

	non_invasive_blood_pressure_diastolic	non_invasive_blood_pressure_systolic
	<dbl>	<dbl>
1	48	84
2	56.5	106
3	102	154
4	90	151
5	93.3	156
6	56	73
7	78	110
8	30.5	174.
9	62	98.5
10	80	112

	respiratory_rate	temperature_fahrenheit
	<dbl>	<dbl>
1	24	98.7
2	24.3	97.7
3	23.5	98
4	18	98.5
5	14	97.6
6	19	97.7
7	16.5	97.9
8	13	98.1
9	14	97.2
10	21	97.9

# i 94,434 more rows

```
# Check for missing values
mimiciv_icu_cohort |> tbl_summary(by = los_long)
```

Characteristic	TRUE N = 46,337 <sup>1</sup>	FALSE N = 48,107 <sup>1</sup>
subject_id	15,021,968 (12,517,625, 17,521,224)	14,988,897 (12,506,011, 17,513,478)
hadm_id	25,011,290 (22,497,215, 27,470,855)	24,954,662 (22,465,369, 27,459,051)
stay_id	34,949,825 (32,473,497, 37,458,915)	35,045,664 (32,534,836, 37,518,493)
gender		
F	20,106 (43%)	21,471 (45%)
M	26,231 (57%)	26,636 (55%)
age_intime	67 (56, 77)	66 (54, 77)
marital_status		
DIVORCED	3,377 (8.0%)	3,555 (8.0%)
MARRIED	20,557 (49%)	21,344 (48%)
SINGLE	12,745 (30%)	14,039 (31%)
WIDOWED	5,319 (13%)	5,752 (13%)
Unknown	4,339	3,417
race		
asian	1,369 (3.0%)	1,516 (3.2%)
black	4,933 (11%)	5,452 (11%)
hispanic	1,687 (3.6%)	1,908 (4.0%)
other	8,036 (17%)	6,880 (14%)
white	30,312 (65%)	32,351 (67%)
first_careunit		

<sup>1</sup> Median (Q1, Q3); n (%)

<b>Characteristic</b>	<b>TRUE</b> N = 46,337 <sup>1</sup>	<b>FALSE</b> N = 48,107 <sup>1</sup>
Cardiac Vascular Intensive Care Unit (CVICU)	7,353 (16%)	7,416 (15%)
Medical Intensive Care Unit (MICU)	9,837 (21%)	10,862 (23%)
Medical/Surgical Intensive Care Unit (MICU/SICU)	6,667 (14%)	8,780 (18%)
Surgical Intensive Care Unit (SICU)	6,434 (14%)	6,574 (14%)
Other	16,046 (35%)	14,475 (30%)
bicarbonate	24.0 (21.0, 27.0)	24.0 (21.0, 27.0)
Unknown	6,272	5,277
chloride	102 (98, 105)	102 (98, 105)
Unknown	6,184	5,167
creatinine	1.00 (0.80, 1.60)	1.00 (0.80, 1.40)
Unknown	4,541	3,486
glucose	122 (100, 159)	118 (98, 154)
Unknown	6,340	5,314
hematocrit	35 (29, 40)	36 (30, 41)
Unknown	3,857	2,894
potassium	4.20 (3.90, 4.70)	4.20 (3.90, 4.60)
Unknown	6,200	5,187
sodium	138.0 (135.0, 141.0)	139.0 (136.0, 141.0)
Unknown	6,167	5,163
wbc	9.7 (7.0, 13.8)	9.0 (6.6, 12.6)
Unknown	3,906	2,944
heart_rate	87 (75, 102)	84 (73, 99)

<sup>1</sup> Median (Q1, Q3); n (%)

Characteristic	TRUE N = 46,337 <sup>1</sup>	FALSE N = 48,107 <sup>1</sup>
Unknown	1	84
non_invasive_blood_pressure_diastolic	67 (57, 79)	68 (58, 80)
Unknown	350	1,015
non_invasive_blood_pressure_systolic	120 (104, 137)	122 (107, 138)
Unknown	347	1,013
respiratory_rate	19.0 (16.0, 23.0)	18.0 (15.0, 22.0)
Unknown	14	181
temperature_fahrenheit	98.20 (97.70, 98.80)	98.10 (97.60, 98.60)
Unknown	230	1,386

<sup>1</sup> Median (Q1, Q3); n (%)

There are missing values in `marital_status`, lab measurements, and vital measurements. Missing values are visualized before deciding how to handle them.

```
numeric_cols <- select(mimiciv_icu_cohort |>
  select(-subject_id, -hadm_id, -stay_id, -los_long,
    -gender, -age_intime, -race, -first_careunit),
  where(is.numeric)) # Select numeric columns

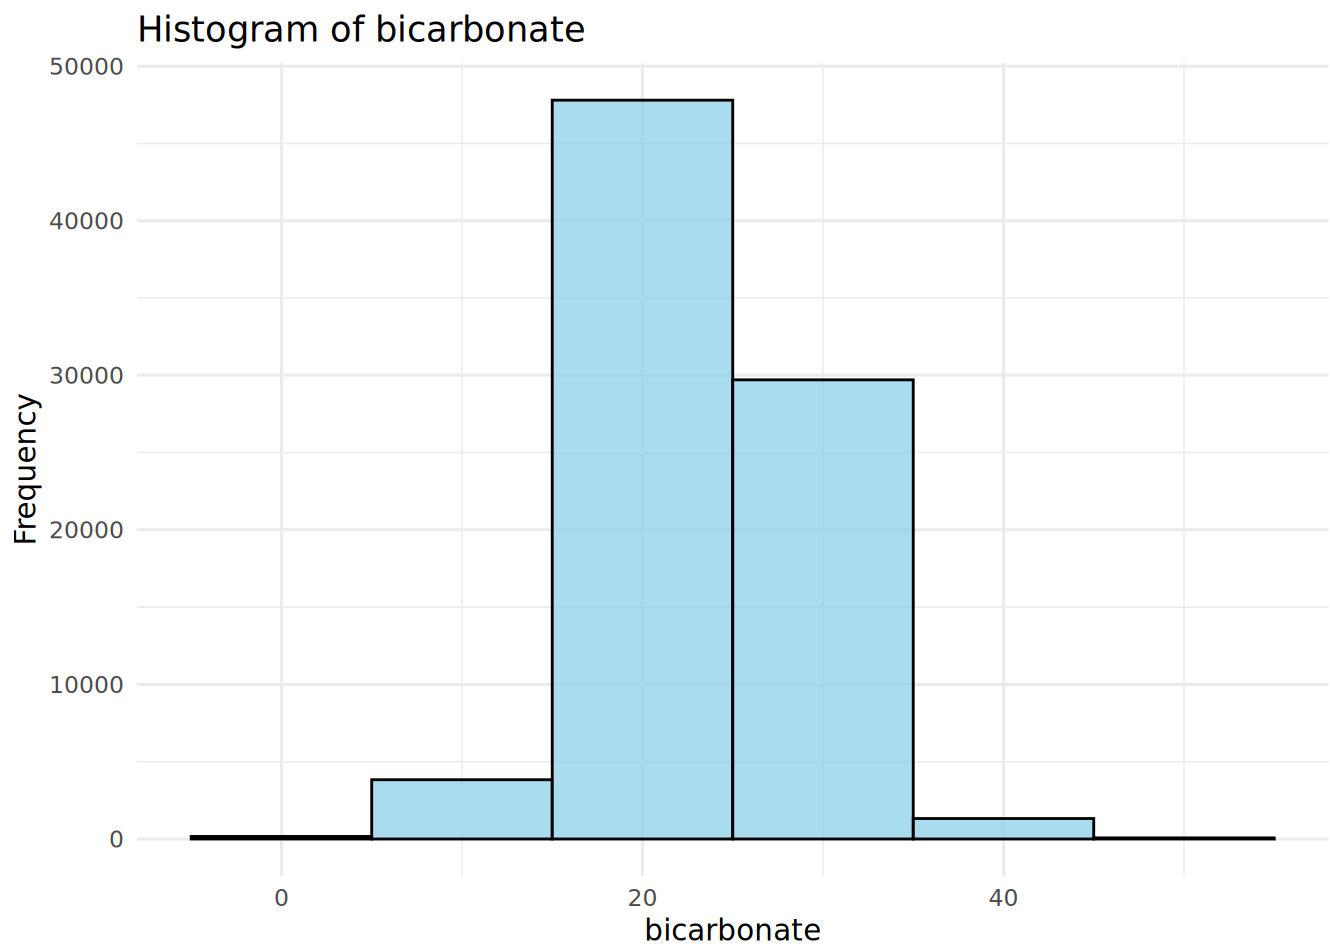
# Loop through each numeric column and create a histogram
for (col in names(numeric_cols)) {
  p <- ggplot(mimiciv_icu_cohort,
    aes_string(x = col)) +
    geom_histogram(binwidth = 10, fill = "skyblue",
      color = "black", alpha = 0.7) +
    labs(title = paste("Histogram of", col), x = col, y = "Frequency") +
    theme_minimal()
  print(p)
}
```

Warning: `aes\_string()` was deprecated in ggplot2 3.0.0.

• Please use tidy evaluation idioms with `aes()`.

• See also `vignette("ggplot2-in-packages")` for more information.

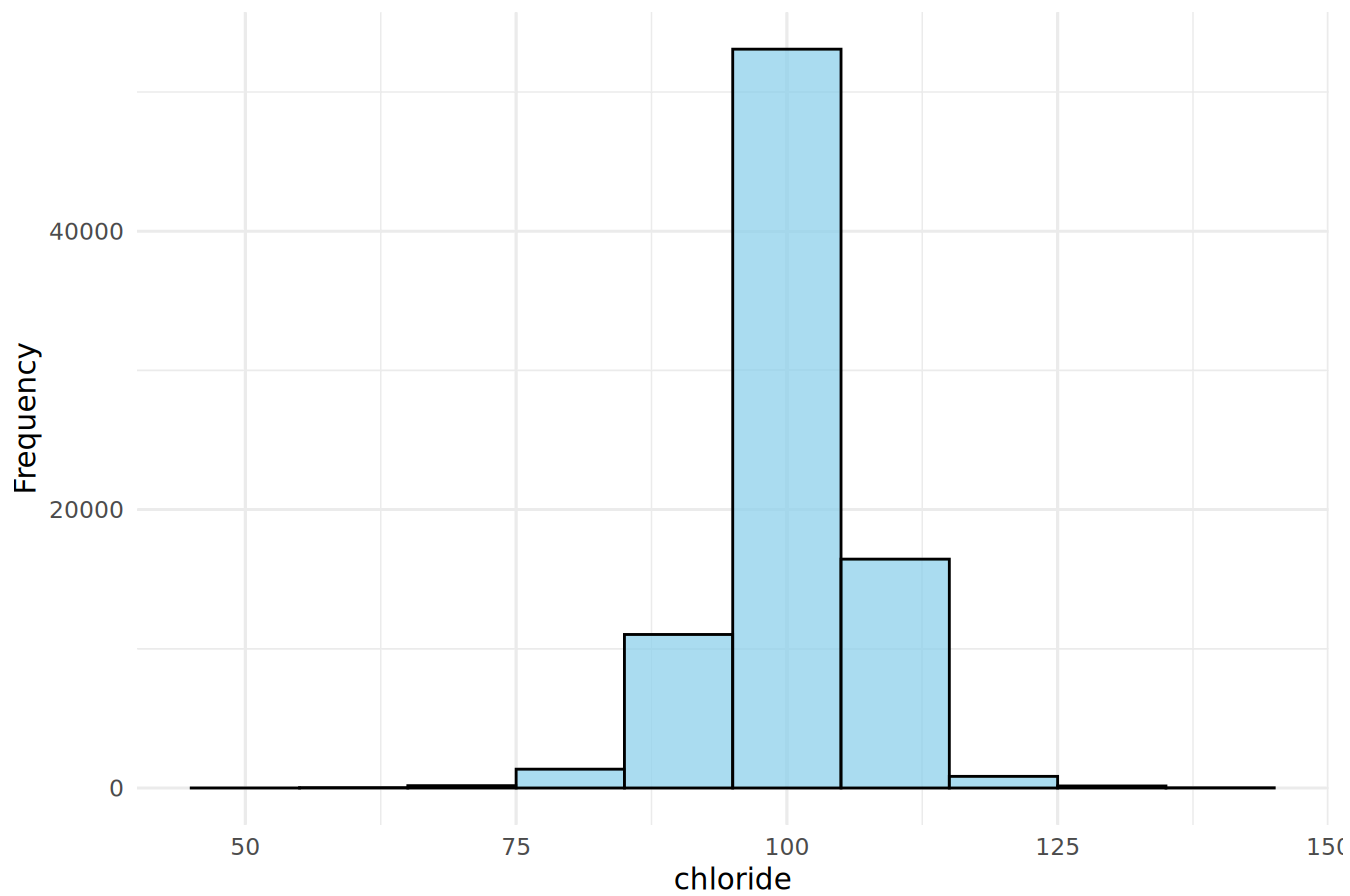
Warning: Removed 11549 rows containing non-finite outside the scale range  
(`stat\_bin()`).



Warning: Removed 11351 rows containing non-finite outside the scale range (``stat_bin()``).

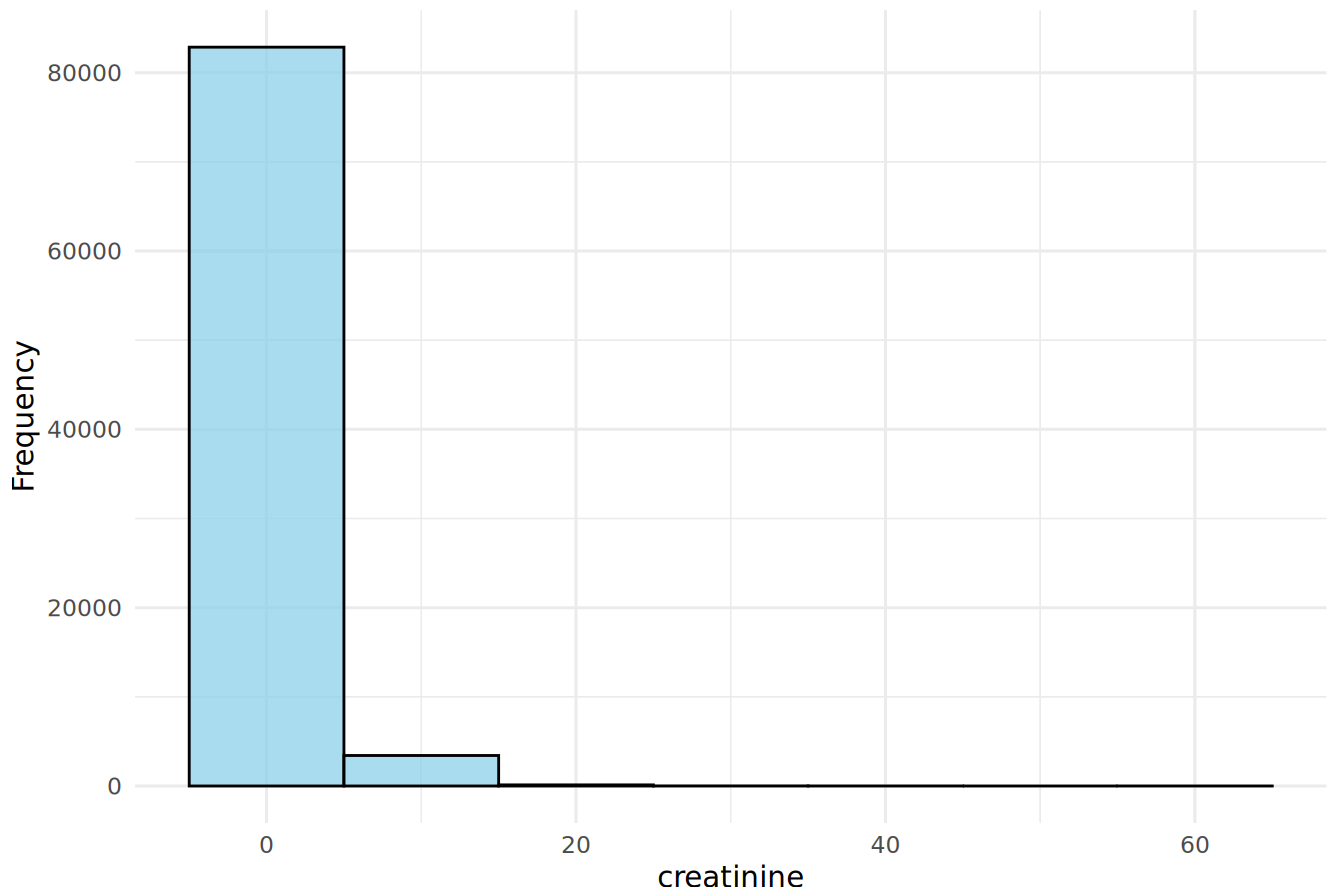


Histogram of chloride



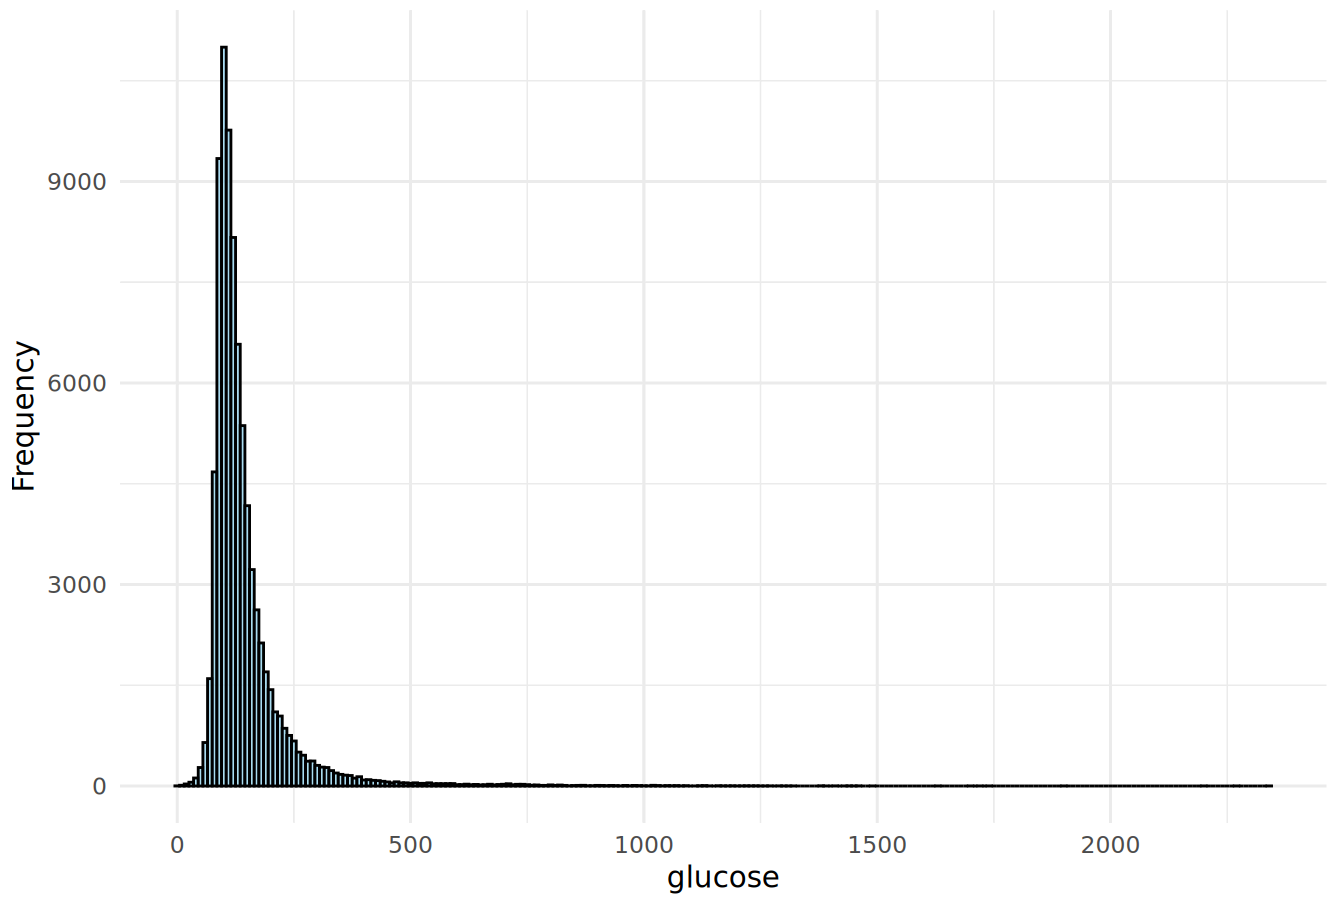
Warning: Removed 8027 rows containing non-finite outside the scale range  
(`stat\_bin()`).

Histogram of creatinine

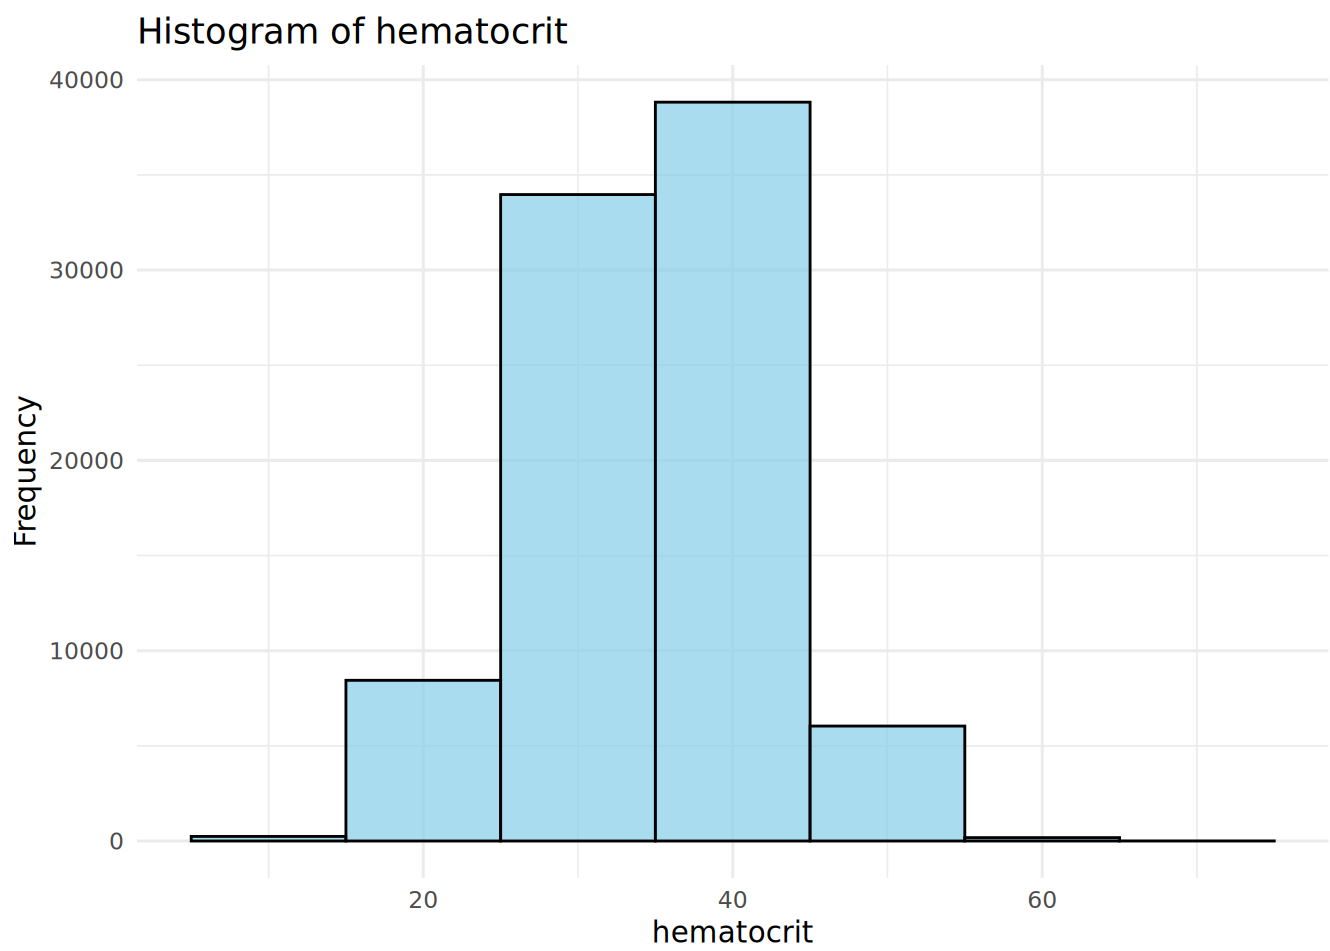


Warning: Removed 11654 rows containing non-finite outside the scale range (``stat_bin()``).

Histogram of glucose

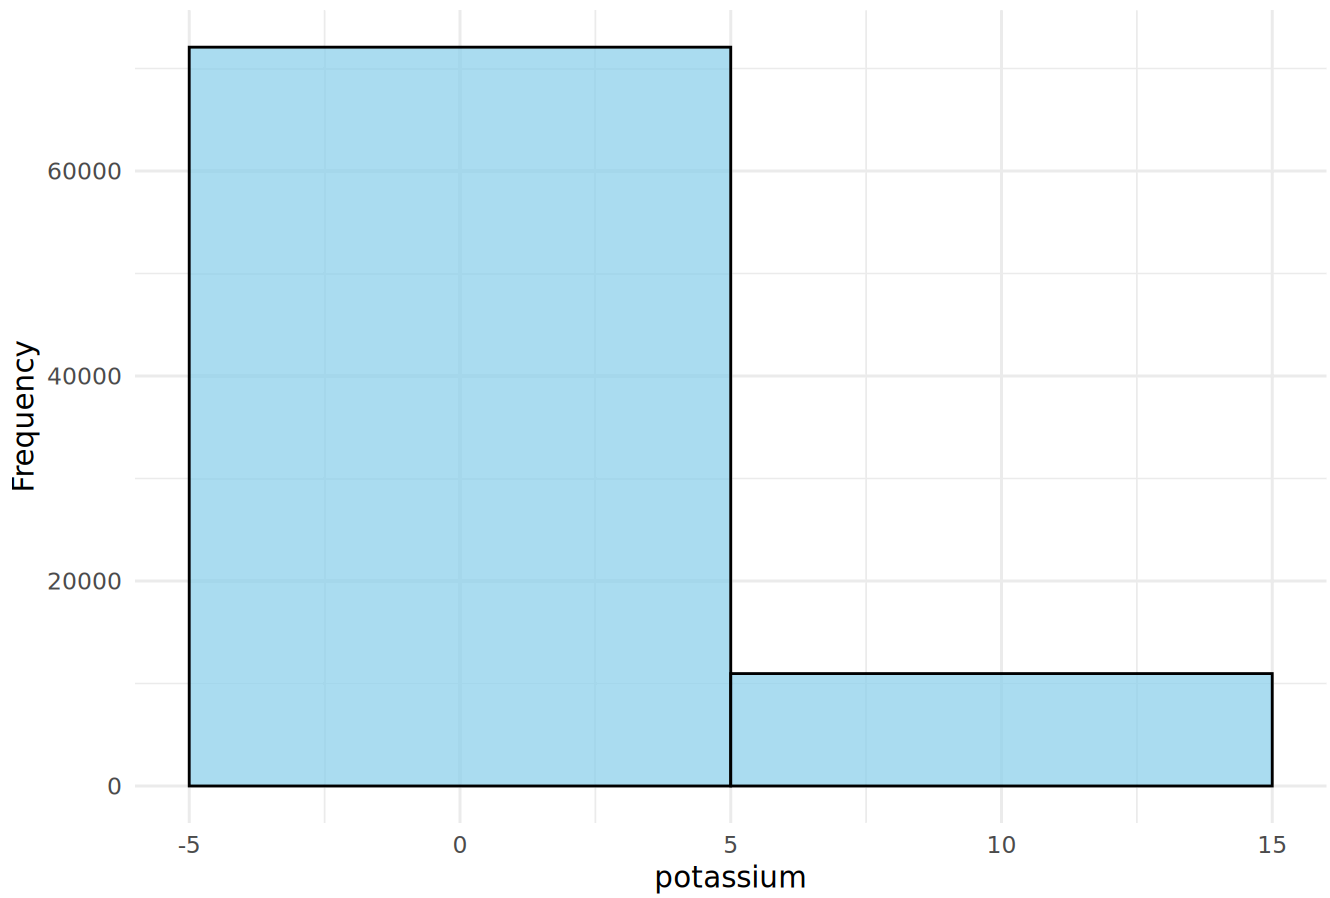


Warning: Removed 6751 rows containing non-finite outside the scale range  
(`stat\_bin()`).

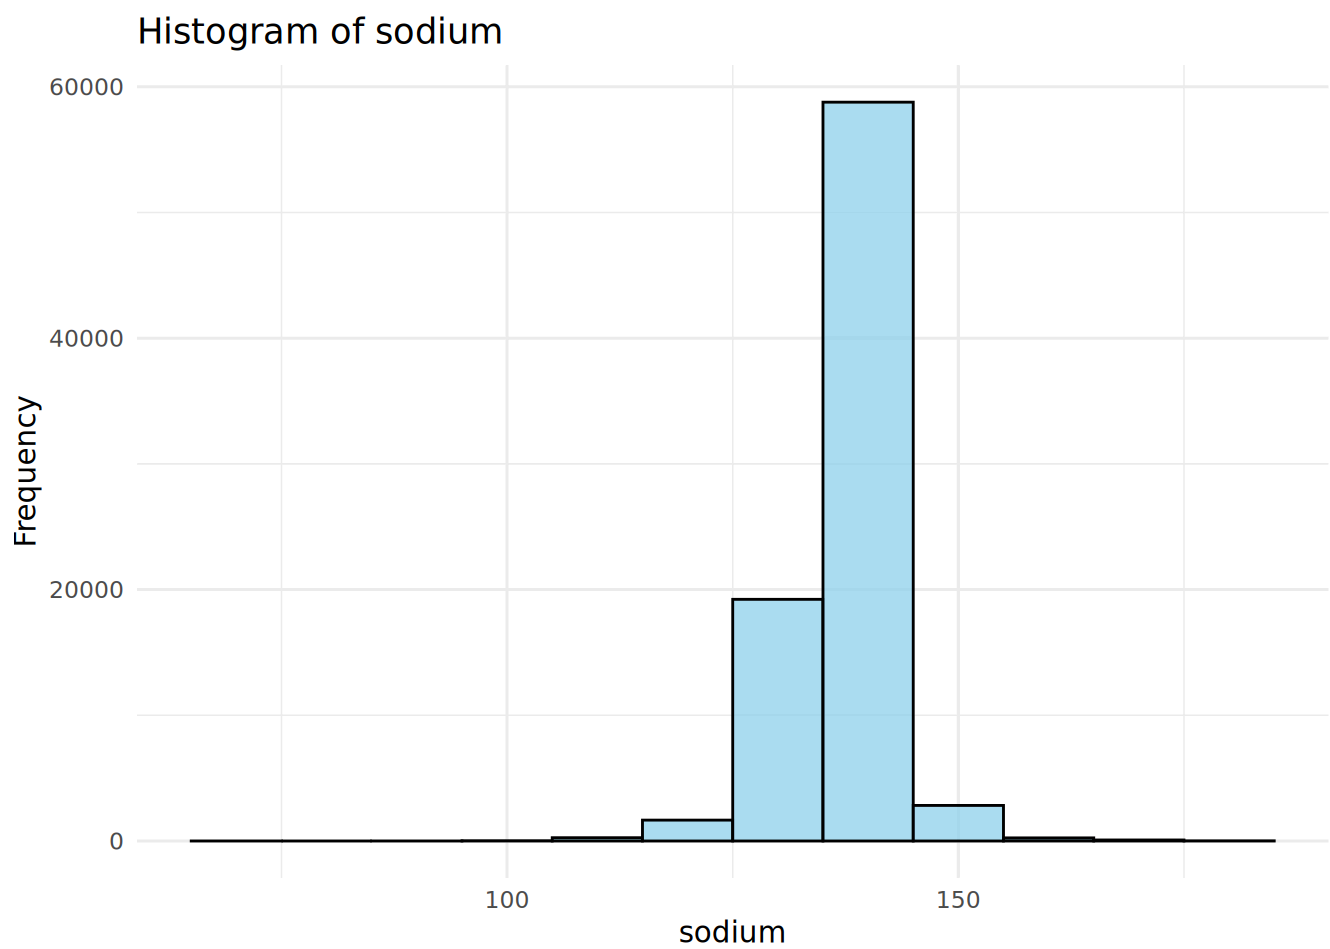


Warning: Removed 11387 rows containing non-finite outside the scale range (``stat_bin()``).

Histogram of potassium

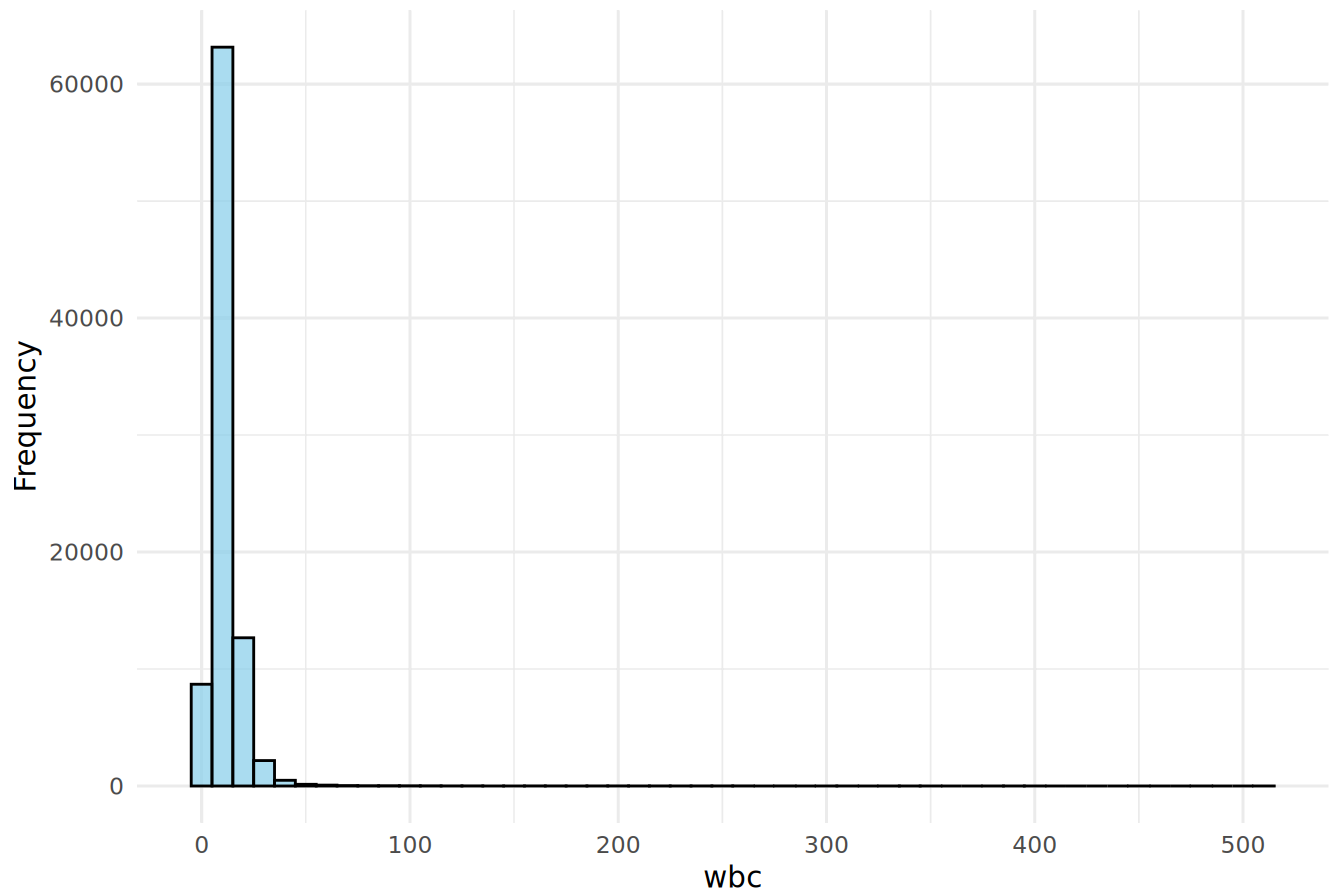


Warning: Removed 11330 rows containing non-finite outside the scale range  
(`stat\_bin()`).



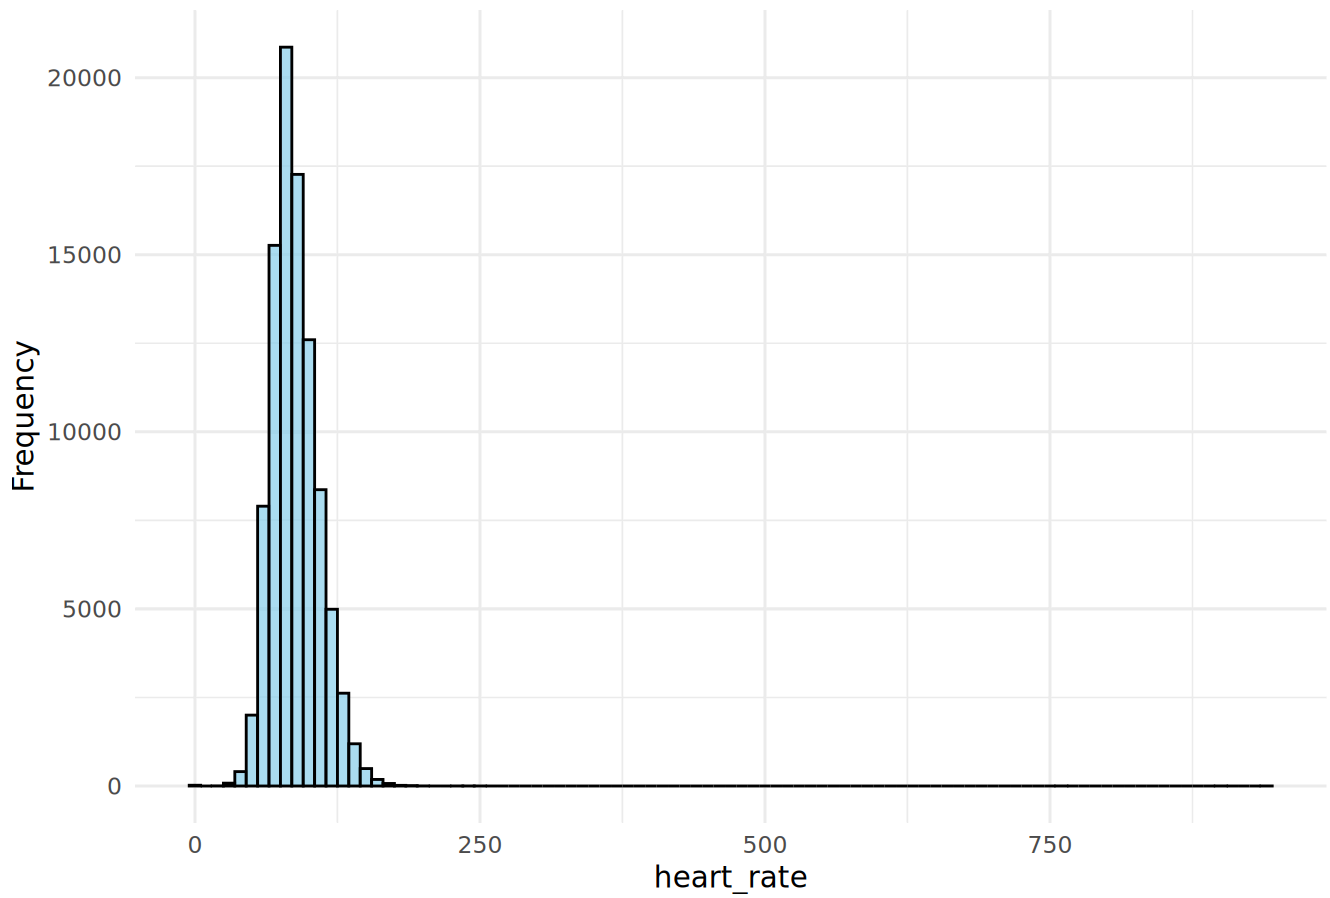
Warning: Removed 6850 rows containing non-finite outside the scale range (``stat_bin()``).

Histogram of wbc



Warning: Removed 85 rows containing non-finite outside the scale range  
(`stat\_bin()`).

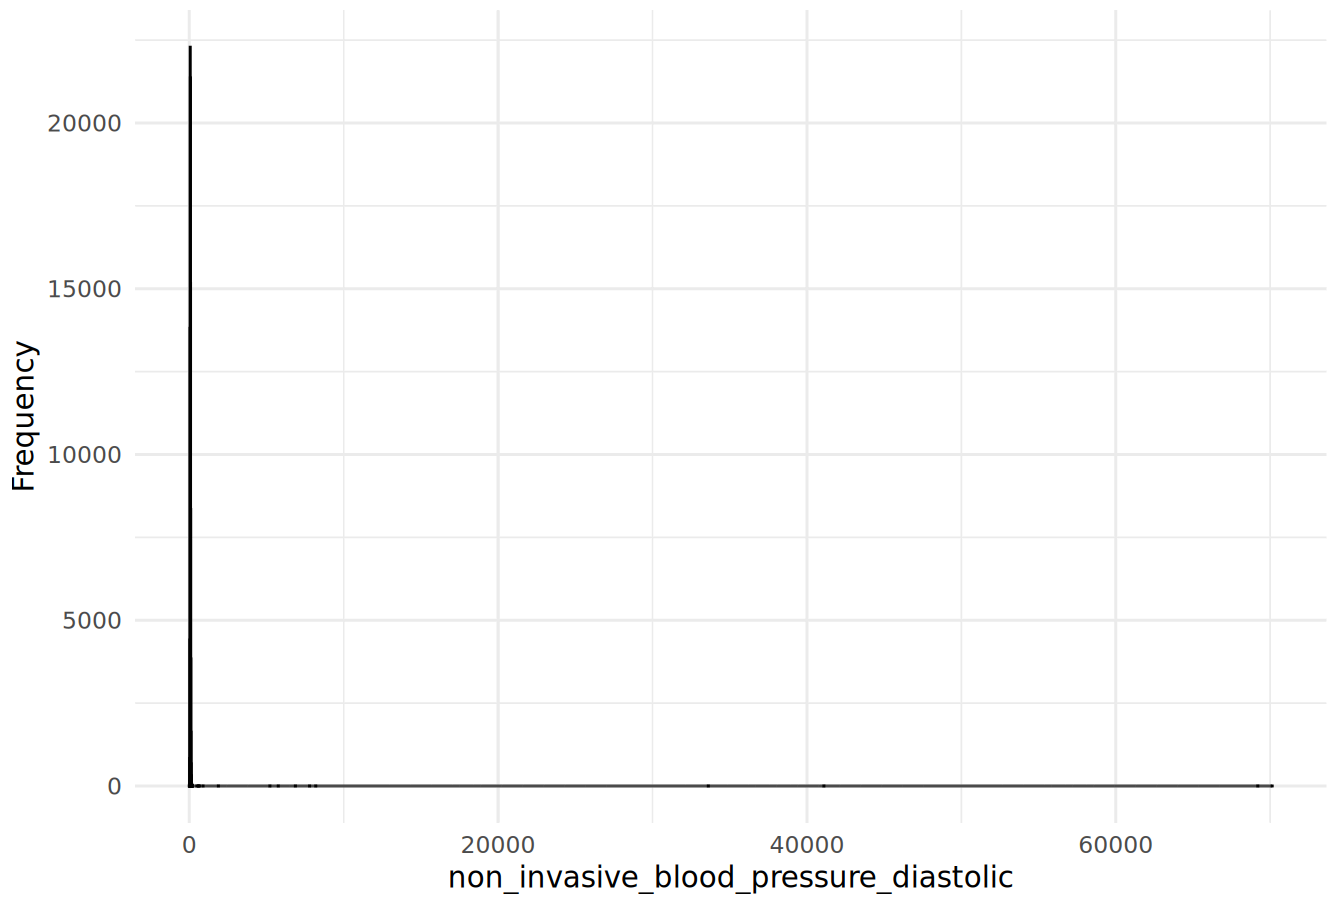
Histogram of heart\_rate



Warning: Removed 1365 rows containing non-finite outside the scale range  
(`stat\_bin()`).

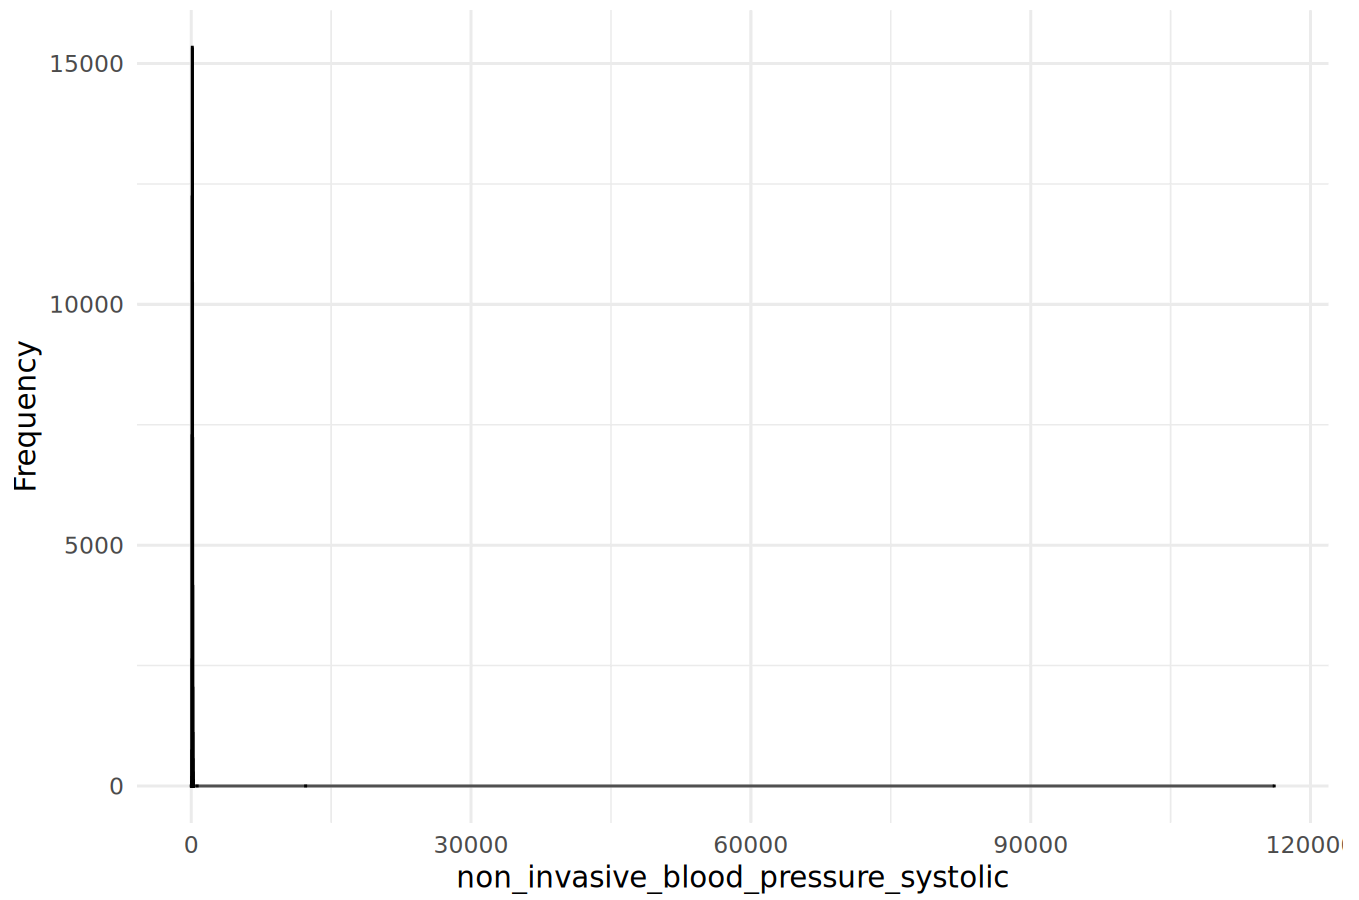


Histogram of non\_invasive\_blood\_pressure\_diastolic

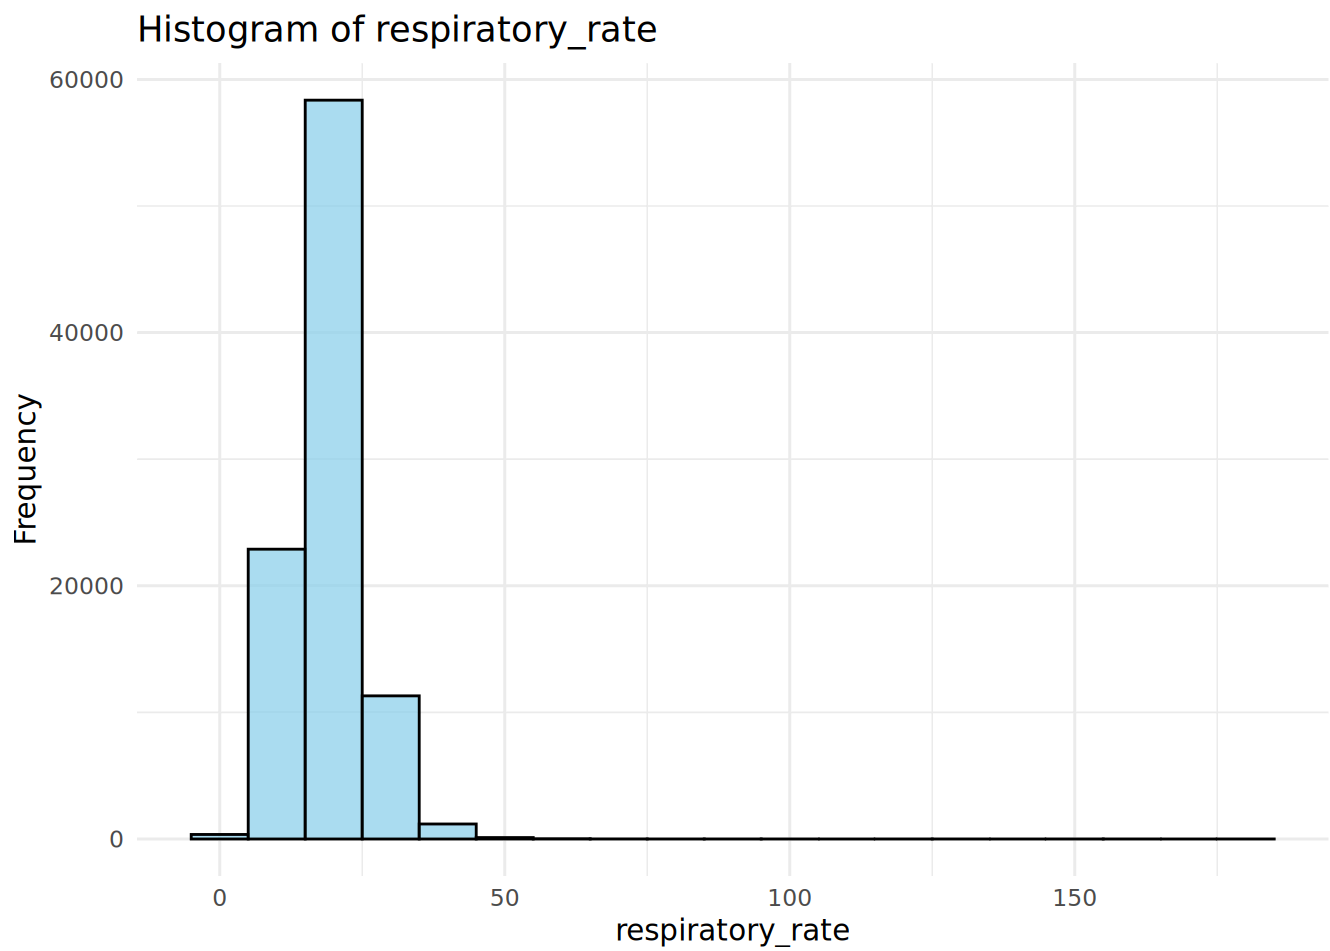


Warning: Removed 1360 rows containing non-finite outside the scale range  
(`stat\_bin()`).

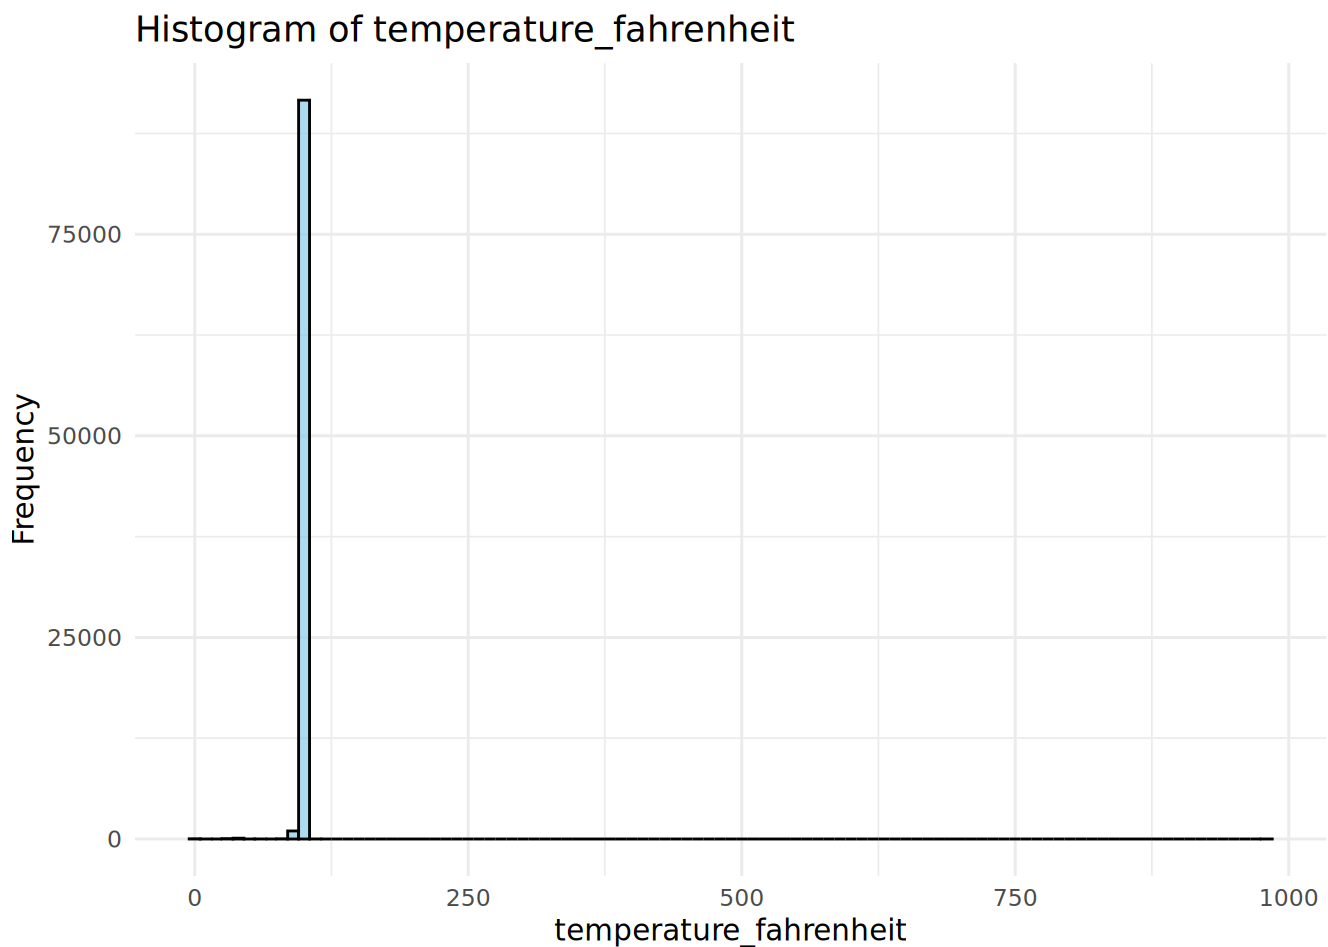
Histogram of non\_invasive\_blood\_pressure\_systolic



Warning: Removed 195 rows containing non-finite outside the scale range  
(`stat\_bin()`).



Warning: Removed 1616 rows containing non-finite outside the scale range (``stat_bin()``).

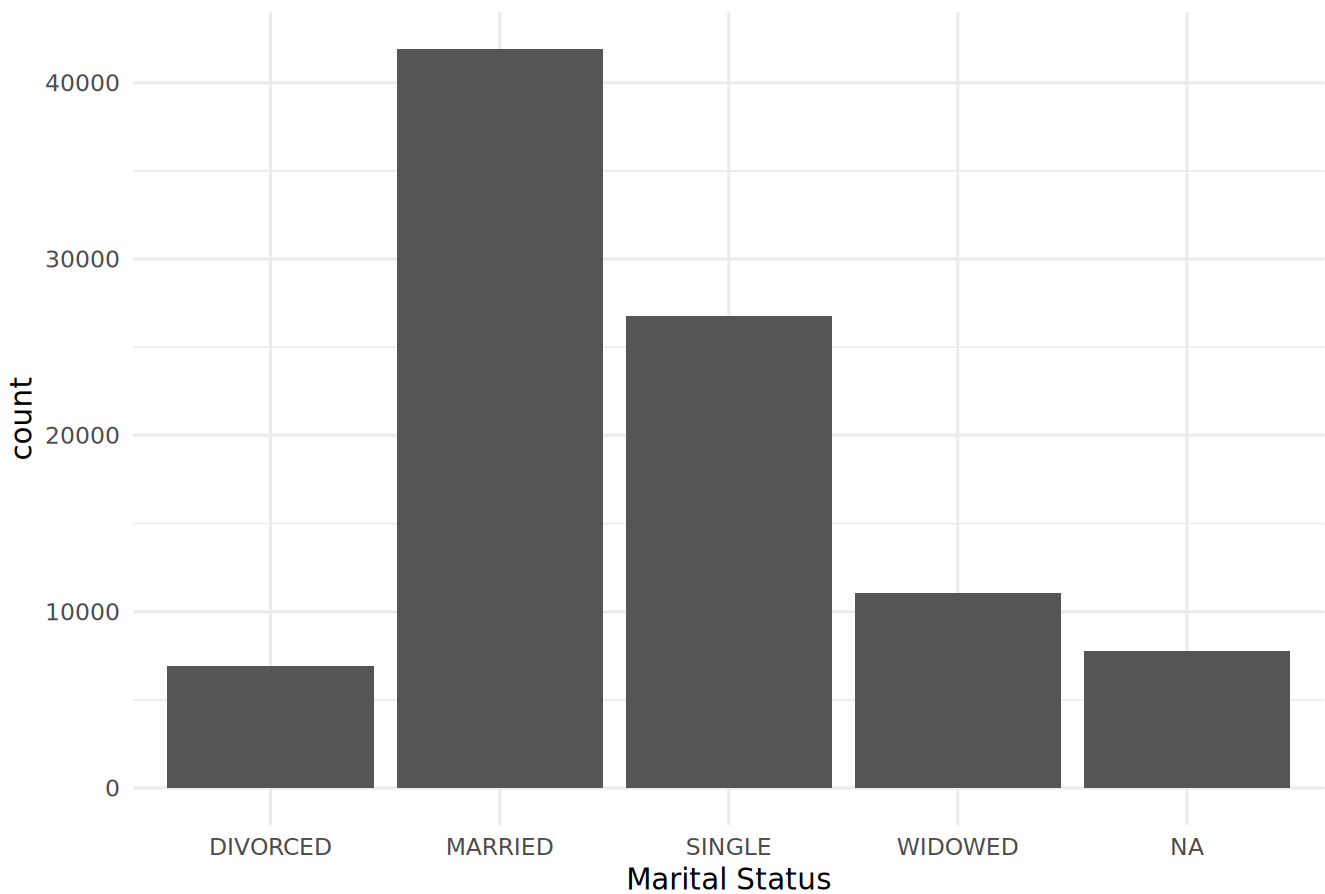


Mean imputation should be used for normally distributed variables, `bicarbonate`, `chloride`, `hematocrit`, `sodium`, and `heart_rate`.

Median imputation should be used for skewed variables, `creatinine`, `glucose`, `potassium`, `wbc`, `non_invasive_blood_pressure_diastolic`, `non_invasive_blood_pressure_systolic`, `respiratory_rate`, and `temperature_fahrenheit`.

```
ggplot(mimiciv_icu_cohort, aes(x = marital_status)) +  
  geom_bar() +  
  labs(title = "Bar plot of Age at ICU Intime by LOS Long",  
        x = "Marital Status") +  
  theme_minimal()
```

Bar plot of Age at ICU Intime by LOS Long



KNN imputation should be used for categorical variables, `marital_status`.

```
rm(cohort_data)
gc()
```

	used (Mb)	gc trigger (Mb)	max used (Mb)
Ncells	2892146 154.5	4883167 260.8	4883167 260.8
Vcells	6748227 51.5	23923775 182.6	23923775 182.6

## 2. Partition data into 50% training set and 50% test set.

Stratify partitioning according to `los_long`. For grading purpose, sort the data by `subject_id`, `hadm_id`, and `stay_id` and use the seed `203` for the initial data split. Below is the sample code.

```
set.seed(203)

# sort
mimiciv_icu_cohort <- mimiciv_icu_cohort |>
  arrange(subject_id, hadm_id, stay_id) |>
  # remove subject_id, hadm_id, stay_id
  select(-subject_id, -hadm_id, -stay_id)

data_split <- initial_split(
```

```

mimiciv_icu_cohort,
# stratify by los_long
strata = "los_long",
prop = 0.5
)

# data_split
mimiciv_icu_cohort_train <- training(data_split)
dim(mimiciv_icu_cohort_train)

```

```
[1] 47221    19
```

```

mimiciv_icu_cohort_test <- testing(data_split)
dim(mimiciv_icu_cohort_test)

```

```
[1] 47223    19
```

```

rm(mimiciv_icu_cohort)
gc()

```

```

           used (Mb) gc trigger (Mb) max used (Mb)
Ncells 2897657 154.8   4883167 260.8  4883167 260.8
Vcells 9806467  74.9   23923775 182.6  23923775 182.6

```

### 3. Train and tune the models using the training set.

Logistic regression with enet regularization

```

logit_rec <- recipe(los_long ~ ., data = mimiciv_icu_cohort_train) |>
# Mean imputation for normal variables
step_impute_mean(bicarbonate, chloride, hematocrit, sodium, heart_rate) |>

# Median imputation for skewed variables
step_impute_median(creatinine, glucose, potassium, wbc,
                    non_invasive_blood_pressure_diastolic,
                    non_invasive_blood_pressure_systolic,
                    respiratory_rate, temperature_fahrenheit) |>

# KNN imputation for categorical variables
step_impute_knn(marital_status) |>

# create traditional dummy variables
step_dummy(all_nominal_predictors()) |>
# zero-variance filter
step_zv(all_nominal_predictors()) |>
# center and scale numeric data
step_normalize(all_numeric_predictors()) |>
print()

```

## — Recipe

---

### — Inputs

Number of variables by role

```
outcome: 1
predictor: 18
```

### — Operations

- Mean imputation for: bicarbonate, chloride, hematocrit, sodium, ...
- Median imputation for: creatinine, glucose, potassium, wbc, ...
- K-nearest neighbor imputation for: marital\_status
- Dummy variables from: all\_nominal\_predictors()
- Zero variance filter on: all\_nominal\_predictors()
- Centering and scaling for: all\_numeric\_predictors()

```
logit_mod <- logistic_reg(
  penalty = tune(),
  mixture = tune()
) |>
  set_engine("glmnet", standardize = FALSE) |>
  print()
```

Logistic Regression Model Specification (classification)

Main Arguments:

```
penalty = tune()
mixture = tune()
```

Engine-Specific Arguments:

```
standardize = FALSE
```

Computational engine: glmnet

```
logit_wf <- workflow() |>
  add_recipe(logit_rec) |>
  add_model(logit_mod) |>
  print()
```

## == Workflow ==

Preprocessor: Recipe

Model: logistic\_reg()

## — Preprocessor —

6 Recipe Steps

- step\_impute\_mean()
- step\_impute\_median()
- step\_impute\_knn()
- step\_dummy()
- step\_zv()
- step\_normalize()

## — Model —

Logistic Regression Model Specification (classification)

Main Arguments:

penalty = tune()

mixture = tune()

Engine-Specific Arguments:

standardize = FALSE

Computational engine: glmnet

```
logit_grid <- grid_regular(  
  penalty(range = c(-4, 1)),  
  mixture(),  
  levels = c(50, 5)  
) |>  
print()
```

# A tibble: 250 × 2

	penalty	mixture
	<dbl>	<dbl>
1	0.0001	0
2	0.000126	0
3	0.000160	0
4	0.000202	0
5	0.000256	0
6	0.000324	0
7	0.000409	0
8	0.000518	0
9	0.000655	0
10	0.000829	0

# i 240 more rows

```
set.seed(203)
```



```

folds <- vfold_cv(mimiciv_icu_cohort_train, v = 5, strata = los_long)
folds

```

```

# 5-fold cross-validation using stratification

```

```

# A tibble: 5 × 2

```

```

  splits      id
  <list>    <chr>
1 <split [37776/9445]> Fold1
2 <split [37776/9445]> Fold2
3 <split [37776/9445]> Fold3
4 <split [37778/9443]> Fold4
5 <split [37778/9443]> Fold5

```

```

logit_fit <- logit_wf |>
  tune_grid(
    resamples = folds,
    grid = logit_grid,
    metrics = metric_set(roc_auc, accuracy)
  )

```

```

logit_fit |>
  # aggregate metrics from K folds
  collect_metrics() |>
  print(width = Inf) |>
  filter(.metric == "roc_auc") |>
  ggplot(mapping = aes(x = penalty, y = mean, color = factor(mixture))) +
  geom_point() +
  labs(x = "Penalty", y = "CV AUC") +
  scale_x_log10()

```

```

# A tibble: 500 × 8

```

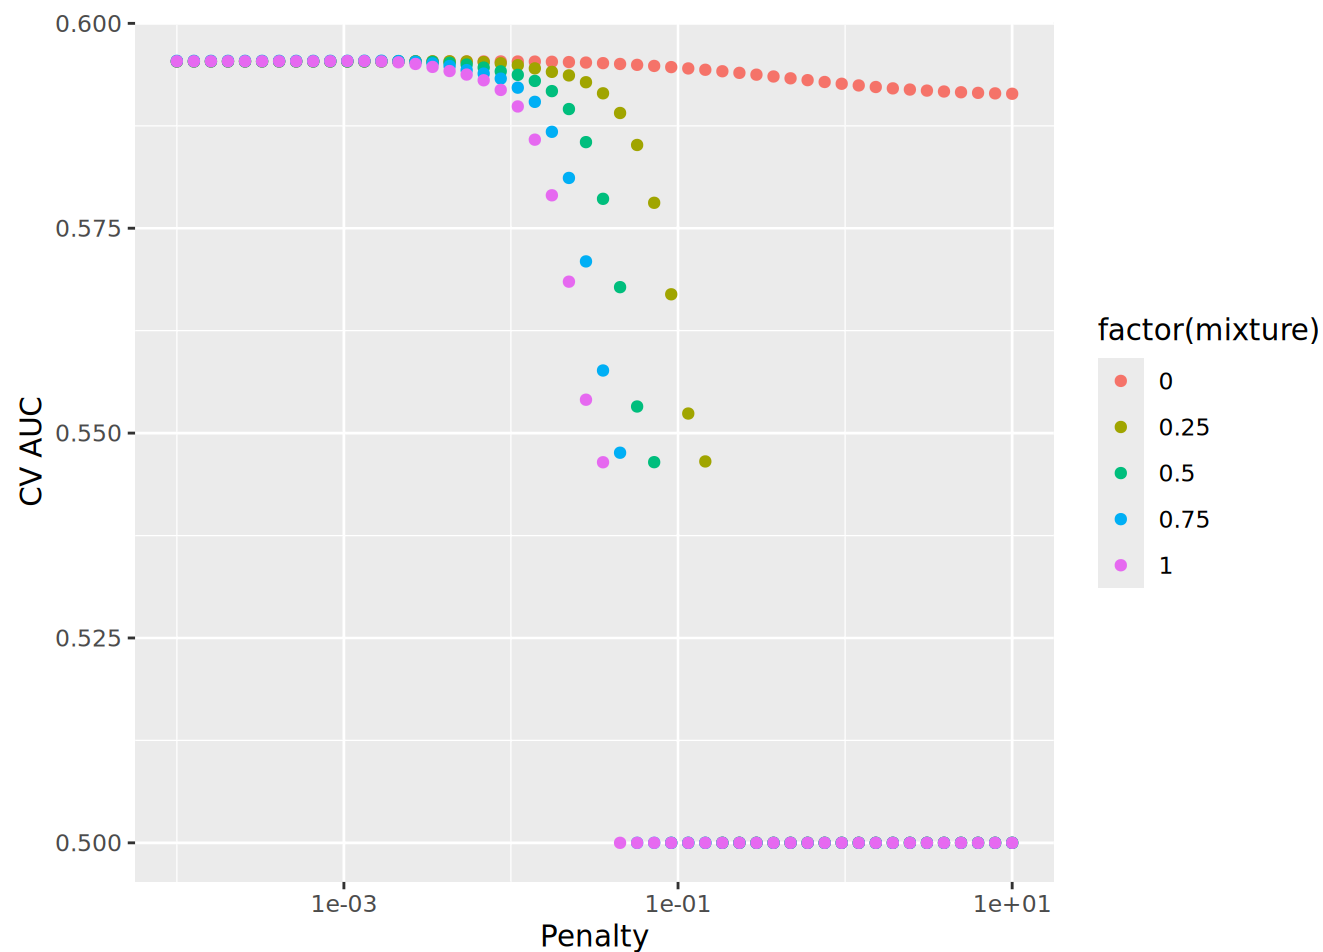
	penalty	mixture	.metric	.estimator	mean	n	std_err
	<dbl>	<dbl>	<chr>	<chr>	<dbl>	<int>	<dbl>
1	0.0001	0	accuracy	binary	0.570	5	0.000757
2	0.0001	0	roc_auc	binary	0.595	5	0.00119
3	0.000126	0	accuracy	binary	0.570	5	0.000757
4	0.000126	0	roc_auc	binary	0.595	5	0.00119
5	0.000160	0	accuracy	binary	0.570	5	0.000757
6	0.000160	0	roc_auc	binary	0.595	5	0.00119
7	0.000202	0	accuracy	binary	0.570	5	0.000757
8	0.000202	0	roc_auc	binary	0.595	5	0.00119
9	0.000256	0	accuracy	binary	0.570	5	0.000757
10	0.000256	0	roc_auc	binary	0.595	5	0.00119
	.config						
	<chr>						
1	Preprocessor1_Model001						
2	Preprocessor1_Model001						
3	Preprocessor1_Model002						
4	Preprocessor1_Model002						

```

5 Preprocessor1_Model1003
6 Preprocessor1_Model1003
7 Preprocessor1_Model1004
8 Preprocessor1_Model1004
9 Preprocessor1_Model1005
10 Preprocessor1_Model1005

```

```
# i 490 more rows
```



```
logit_fit |> show_best(metric = "roc_auc")
```

```
# A tibble: 5 × 8
```

	penalty	mixture	.metric	.estimator	mean	n	std_err	.config
	<dbl>	<dbl>	<chr>	<chr>	<dbl>	<int>	<dbl>	<chr>
1	0.00168	0.75	roc_auc	binary	0.595	5	0.00113	Preprocessor1_Model163
2	0.00133	0.75	roc_auc	binary	0.595	5	0.00115	Preprocessor1_Model162
3	0.00105	0.75	roc_auc	binary	0.595	5	0.00117	Preprocessor1_Model161
4	0.000829	0.75	roc_auc	binary	0.595	5	0.00120	Preprocessor1_Model160
5	0.0001	0.75	roc_auc	binary	0.595	5	0.00120	Preprocessor1_Model151

```
logit_best <- logit_fit |>
  select_best(metric = "roc_auc")
```

```
logit_best
```

```
# A tibble: 1 × 3
  penalty mixture .config
  <dbl>   <dbl> <chr>
1 0.00168    0.75 Preprocessor1_Model163
```

```
logit_final <- logit_wf |>
  finalize_workflow(logit_best)
logit_final
```

## == Workflow ==

Preprocessor: Recipe

Model: logistic\_reg()

## — Preprocessor —

6 Recipe Steps

- step\_impute\_mean()
- step\_impute\_median()
- step\_impute\_knn()
- step\_dummy()
- step\_zv()
- step\_normalize()

## — Model —

Logistic Regression Model Specification (classification)

Main Arguments:

```
penalty = 0.00167683293681101
mixture = 0.75
```

Engine-Specific Arguments:

```
standardize = FALSE
```

Computational engine: glmnet

```
logit_fit_final <- logit_final |>
  last_fit(data_split)
```

## Random Forest

```
rf_rec <- recipe(los_long ~ ., data = mimiv_icu_cohort_train) |>
  # Mean imputation for normal variables
  step_impute_mean(bicarbonate, chloride, hematocrit, sodium, heart_rate) |>

  # Median imputation for skewed variables
  step_impute_median(creatinine, glucose, potassium, wbc,
                    non_invasive_blood_pressure_diastolic,
                    non_invasive_blood_pressure_systolic,
                    respiratory_rate, temperature_fahrenheit) |>
```

```
# KNN imputation for categorical variables
step_impute_knn(marital_status) |>

# zero-variance filter
step_zv(all_nominal_predictors()) |>
print()
```

## — Recipe

---

### — Inputs

Number of variables by role

```
outcome:    1
predictor: 18
```

### — Operations

- Mean imputation for: bicarbonate, chloride, hematocrit, sodium, ...
- Median imputation for: creatinine, glucose, potassium, wbc, ...
- K-nearest neighbor imputation for: marital\_status
- Zero variance filter on: all\_nominal\_predictors()

```
rf_mod <- rand_forest(
  mode = "classification",
  mtry = tune(), # number of predictors randomly sampled in each split
  trees = tune() # number of trees in ensemble
) |>
  set_engine("ranger")
rf_mod
```

### Random Forest Model Specification (classification)

Main Arguments:

```
mtry = tune()
trees = tune()
```

Computational engine: ranger

```
rf_wf <- workflow() |>
  add_recipe(rf_rec) |>
  add_model(rf_mod)
rf_wf
```

## == Workflow ==

Preprocessor: Recipe

Model: rand\_forest()

## — Preprocessor —

### 4 Recipe Steps

- step\_impute\_mean()
- step\_impute\_median()
- step\_impute\_knn()
- step\_zv()

## — Model —

Random Forest Model Specification (classification)

Main Arguments:

```
mtry = tune()
trees = tune()
```

Computational engine: ranger

```
rf_grid <- grid_regular(
  trees(range = c(200L, 800L)),
  mtry(range = c(1L, 8L)),
  levels = c(2, 4)
)
rf_grid
```

# A tibble: 8 × 2

	trees	mtry
	<int>	<int>
1	200	1
2	800	1
3	200	3
4	800	3
5	200	5
6	800	5
7	200	8
8	800	8

```
set.seed(203)
```

```
folds <- vfold_cv(mimiciv_icu_cohort_train, v = 5, strata = los_long)
folds
```

# 5-fold cross-validation using stratification

# A tibble: 5 × 2

splits	id
<list>	<chr>

```

1 <split [37776/9445]> Fold1
2 <split [37776/9445]> Fold2
3 <split [37776/9445]> Fold3
4 <split [37778/9443]> Fold4
5 <split [37778/9443]> Fold5

```

```

rf_fit <- rf_wf |>
  tune_grid(
    resamples = folds,
    grid = rf_grid,
    metrics = metric_set(roc_auc, accuracy)
  )
rf_fit

```

```

# Tuning results
# 5-fold cross-validation using stratification
# A tibble: 5 × 4

```

splits	id	.metrics	.notes
<list>	<chr>	<list>	<list>
1 <split [37776/9445]>	Fold1	<tibble [40 × 6]>	<tibble [0 × 3]>
2 <split [37776/9445]>	Fold2	<tibble [40 × 6]>	<tibble [0 × 3]>
3 <split [37776/9445]>	Fold3	<tibble [40 × 6]>	<tibble [0 × 3]>
4 <split [37778/9443]>	Fold4	<tibble [40 × 6]>	<tibble [0 × 3]>
5 <split [37778/9443]>	Fold5	<tibble [40 × 6]>	<tibble [0 × 3]>

```

rf_fit |>
  collect_metrics() |>
  print(width = Inf) |>
  filter(.metric == "roc_auc") |>
  ggplot(mapping = aes(x = trees, y = mean, color = factor(mtry))) +
  geom_point() +
  labs(x = "Num. of Trees", y = "CV AUC")

```

```

# A tibble: 40 × 8

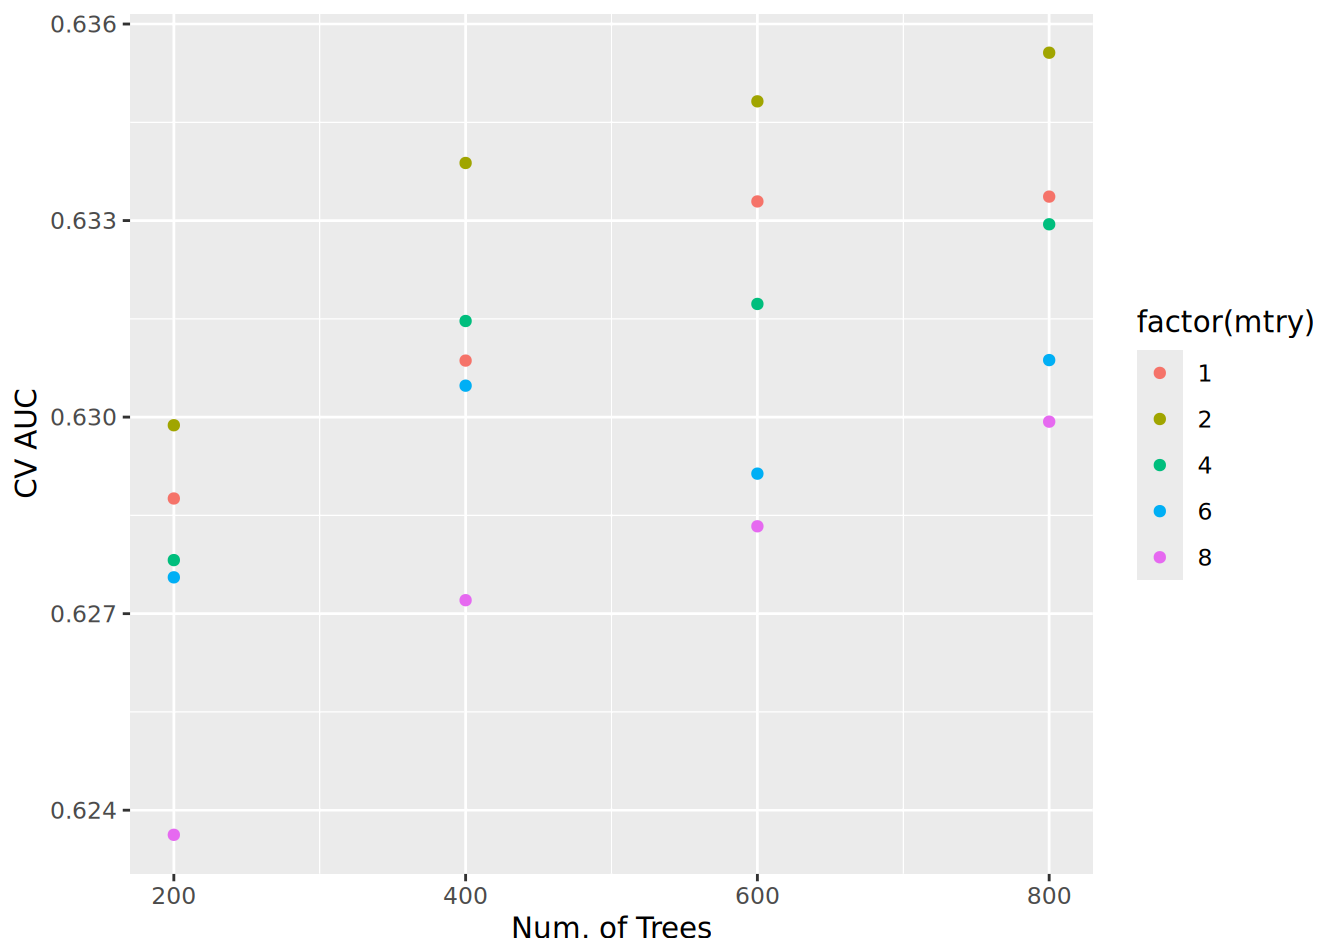
```

	mtry	trees	.metric	.estimator	mean	n	std_err	.config
	<int>	<int>	<chr>	<chr>	<dbl>	<int>	<dbl>	<chr>
1	1	200	accuracy	binary	0.592	5	0.00269	Preprocessor1_Model01
2	1	200	roc_auc	binary	0.629	5	0.00281	Preprocessor1_Model01
3	1	400	accuracy	binary	0.591	5	0.00225	Preprocessor1_Model02
4	1	400	roc_auc	binary	0.631	5	0.00207	Preprocessor1_Model02
5	1	600	accuracy	binary	0.595	5	0.00270	Preprocessor1_Model03
6	1	600	roc_auc	binary	0.633	5	0.00246	Preprocessor1_Model03
7	1	800	accuracy	binary	0.596	5	0.00232	Preprocessor1_Model04
8	1	800	roc_auc	binary	0.633	5	0.00243	Preprocessor1_Model04
9	2	200	accuracy	binary	0.593	5	0.00274	Preprocessor1_Model05
10	2	200	roc_auc	binary	0.630	5	0.00248	Preprocessor1_Model05

```

# i 30 more rows

```



```
rf_fit |> show_best(metric = "roc_auc")
```

# A tibble: 5 × 8

	mtry	trees	.metric	.estimator	mean	n	std_err	.config
	<int>	<int>	<chr>	<chr>	<dbl>	<int>	<dbl>	<chr>
1	2	800	roc_auc	binary	0.636	5	0.00220	Preprocessor1_Model08
2	2	600	roc_auc	binary	0.635	5	0.00210	Preprocessor1_Model07
3	2	400	roc_auc	binary	0.634	5	0.00266	Preprocessor1_Model06
4	1	800	roc_auc	binary	0.633	5	0.00243	Preprocessor1_Model04
5	1	600	roc_auc	binary	0.633	5	0.00246	Preprocessor1_Model03

```
rf_best <- rf_fit |>
  select_best(metric = "roc_auc")
rf_best
```

# A tibble: 1 × 3

	mtry	trees	.config
	<int>	<int>	<chr>
1	2	800	Preprocessor1_Model08

```
rf_final <- rf_wf |>
  finalize_workflow(rf_best)
```

```
rf_final
```

---

## == Workflow ==

Preprocessor: Recipe

Model: rand\_forest()

---

## — Preprocessor —

### 4 Recipe Steps

- step\_impute\_mean()
- step\_impute\_median()
- step\_impute\_knn()
- step\_zv()

---

## — Model —

Random Forest Model Specification (classification)

Main Arguments:

mtry = 2

trees = 800

Computational engine: ranger

```
rf_fit_final <- rf_final |>
  last_fit(data_split)
rf_fit_final
```

# Resampling results

# Manual resampling

# A tibble: 1 × 6

splits	id	.metrics	.notes	.predictions	.workflow
<list>	<chr>	<list>	<list>	<list>	<list>
1 <split [47221/47223]>	train/test sp...	<tibble>	<tibble>	<tibble>	<workflow>

## XGBoost

```
gb_rec <- recipe(los_long ~ ., data = mimiciu_icu_cohort_train) |>
  # Mean imputation for normal variables
  step_impute_mean(bicarbonate, chloride, hematocrit, sodium, heart_rate) |>

  # Median imputation for skewed variables
  step_impute_median(creatinine, glucose, potassium, wbc,
                     non_invasive_blood_pressure_diastolic,
                     non_invasive_blood_pressure_systolic,
                     respiratory_rate, temperature_fahrenheit) |>

  # KNN imputation for categorical variables
  step_impute_knn(marital_status) |>
```



```
# create dummy variable
step_dummy(all_nominal_predictors()) |>

# zero-variance filter
step_zv(all_nominal_predictors()) |>
print()
```

---

## — Recipe

### — Inputs

Number of variables by role

outcome: 1  
predictor: 18

### — Operations

- Mean imputation for: bicarbonate, chloride, hematocrit, sodium, ...
- Median imputation for: creatinine, glucose, potassium, wbc, ...
- K-nearest neighbor imputation for: marital\_status
- Dummy variables from: all\_nominal\_predictors()
- Zero variance filter on: all\_nominal\_predictors()

```
gb_mod <- boost_tree(  
  mode = "classification",  
  trees = 1000,  
  tree_depth = tune(),  
  learn_rate = tune()  
) |>  
  set_engine("xgboost")  
gb_mod
```

## Boosted Tree Model Specification (classification)

### Main Arguments:

```
trees = 1000  
tree_depth = tune()  
learn_rate = tune()
```

Computational engine: xgboost

```
gb_wf <- workflow() |>
  add_recipe(gb_rec) |>
  add_model(gb_mod)
gb_wf
```

## == Workflow ==

Preprocessor: Recipe

Model: boost\_tree()

## — Preprocessor —

5 Recipe Steps

- step\_impute\_mean()
- step\_impute\_median()
- step\_impute\_knn()
- step\_dummy()
- step\_zv()

## — Model —

Boosted Tree Model Specification (classification)

Main Arguments:

trees = 1000

tree\_depth = tune()

learn\_rate = tune()

Computational engine: xgboost

```
gb_grid <- grid_regular(
  tree_depth(range = c(3L, 10L)),
  learn_rate(range = c(0.01, 0.3)),
  levels = c(2, 2)
)
gb_grid
```

# A tibble: 4 × 2

	tree_depth	learn_rate
	<int>	<dbl>
1	3	1.02
2	10	1.02
3	3	2.00
4	10	2.00

```
set.seed(203)
```

```
folds <- vfold_cv(mimiciv_icu_cohort_train, v = 5, strata = los_long)
folds
```

```
# 5-fold cross-validation using stratification
# A tibble: 5 × 2
  splits          id
  <list>         <chr>
1 <split [37776/9445]> Fold1
2 <split [37776/9445]> Fold2
3 <split [37776/9445]> Fold3
4 <split [37778/9443]> Fold4
5 <split [37778/9443]> Fold5
```

```
gb_fit <- gb_wf |>
  tune_grid(
    resamples = folds,
    grid = gb_grid,
    metrics = metric_set(roc_auc, accuracy)
  )
gb_fit
```

```
# Tuning results
# 5-fold cross-validation using stratification
# A tibble: 5 × 4
  splits          id  .metrics          .notes
  <list>         <chr> <list>          <list>
1 <split [37776/9445]> Fold1 <tibble [20 × 6]> <tibble [0 × 3]>
2 <split [37776/9445]> Fold2 <tibble [20 × 6]> <tibble [0 × 3]>
3 <split [37776/9445]> Fold3 <tibble [20 × 6]> <tibble [0 × 3]>
4 <split [37778/9443]> Fold4 <tibble [20 × 6]> <tibble [0 × 3]>
5 <split [37778/9443]> Fold5 <tibble [20 × 6]> <tibble [0 × 3]>
```

```
gb_fit |>
  collect_metrics() |>
  print(width = Inf) |>
  filter(.metric == "roc_auc") |>
  ggplot(mapping = aes(x = learn_rate, y = mean, color = factor(tree_depth))) +
  geom_point() +
  labs(x = "Learning Rate", y = "CV AUC")
```

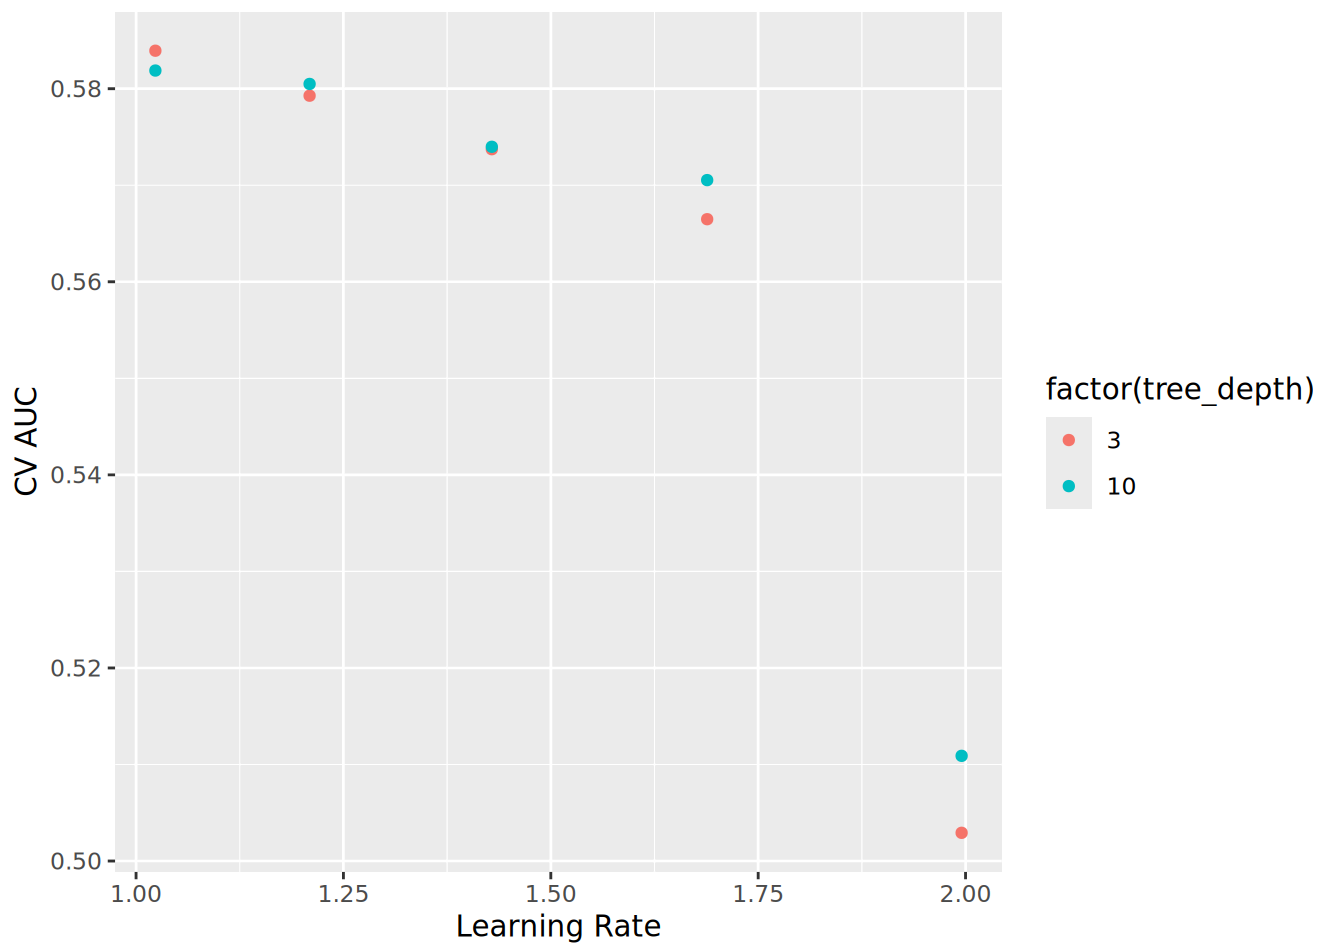
```
# A tibble: 20 × 8
  tree_depth learn_rate .metric .estimator mean    n std_err
    <int>      <dbl> <chr>    <chr>    <dbl> <int>  <dbl>
1         3      1.02 accuracy binary    0.560     5 0.00207
2         3      1.02 roc_auc   binary    0.584     5 0.00245
3        10      1.02 accuracy binary    0.560     5 0.00165
4        10      1.02 roc_auc   binary    0.582     5 0.00135
5         3      1.21 accuracy binary    0.559     5 0.00146
6         3      1.21 roc_auc   binary    0.579     5 0.00180
7        10      1.21 accuracy binary    0.557     5 0.00105
8        10      1.21 roc_auc   binary    0.580     5 0.00115
9         3      1.43 accuracy binary    0.556     5 0.00158
```

10	3	1.43	roc_auc	binary	0.574	5 0.00120
11	10	1.43	accuracy	binary	0.550	5 0.00153
12	10	1.43	roc_auc	binary	0.574	5 0.000296
13	3	1.69	accuracy	binary	0.552	5 0.00186
14	3	1.69	roc_auc	binary	0.566	5 0.00200
15	10	1.69	accuracy	binary	0.549	5 0.00100
16	10	1.69	roc_auc	binary	0.571	5 0.000592
17	3	2.00	accuracy	binary	0.502	5 0.0103
18	3	2.00	roc_auc	binary	0.503	5 0.00766
19	10	2.00	accuracy	binary	0.512	5 0.00688
20	10	2.00	roc_auc	binary	0.511	5 0.00669

.config

<chr>

1 Preprocessor1\_Model01  
2 Preprocessor1\_Model01  
3 Preprocessor1\_Model02  
4 Preprocessor1\_Model02  
5 Preprocessor1\_Model03  
6 Preprocessor1\_Model03  
7 Preprocessor1\_Model04  
8 Preprocessor1\_Model04  
9 Preprocessor1\_Model05  
10 Preprocessor1\_Model05  
11 Preprocessor1\_Model06  
12 Preprocessor1\_Model06  
13 Preprocessor1\_Model07  
14 Preprocessor1\_Model07  
15 Preprocessor1\_Model08  
16 Preprocessor1\_Model08  
17 Preprocessor1\_Model09  
18 Preprocessor1\_Model09  
19 Preprocessor1\_Model10  
20 Preprocessor1\_Model10



```
gb_fit |> show_best(metric = "roc_auc")
```

# A tibble: 5 × 8

	tree_depth	learn_rate	.metric	.estimator	mean	n	std_err	.config
	<int>	<dbl>	<chr>	<chr>	<dbl>	<int>	<dbl>	<chr>
1	3	1.02	roc_auc	binary	0.584	5	0.00245	Preprocessor1_M...
2	10	1.02	roc_auc	binary	0.582	5	0.00135	Preprocessor1_M...
3	10	1.21	roc_auc	binary	0.580	5	0.00115	Preprocessor1_M...
4	3	1.21	roc_auc	binary	0.579	5	0.00180	Preprocessor1_M...
5	10	1.43	roc_auc	binary	0.574	5	0.000296	Preprocessor1_M...

```
gb_best <- gb_fit |>
  select_best(metric = "roc_auc")
gb_best
```

# A tibble: 1 × 3

	tree_depth	learn_rate	.config
	<int>	<dbl>	<chr>
1	3	1.02	Preprocessor1_Model01

```
gb_final <- gb_wf |>
  finalize_workflow(gb_best)
```

```
gb_final
```

---

## == Workflow ==

Preprocessor: Recipe

Model: boost\_tree()

---

## — Preprocessor —

5 Recipe Steps

- step\_impute\_mean()
- step\_impute\_median()
- step\_impute\_knn()
- step\_dummy()
- step\_zv()

---

## — Model —

Boosted Tree Model Specification (classification)

Main Arguments:

```
trees = 1000  
tree_depth = 3  
learn_rate = 1.02329299228075
```

Computational engine: xgboost

```
gb_fit_final <- gb_final |>  
  last_fit(data_split)  
gb_fit_final
```

# Resampling results

# Manual resampling

# A tibble: 1 × 6

splits	id	.metrics	.notes	.predictions	.workflow
<list>	<chr>	<list>	<list>	<list>	<list>
1 <split [47221/47223]>	train/test sp...	<tibble>	<tibble>	<tibble>	<workflow>

## Model Stacking

```
stacks_recipe <- recipe(los_long ~ ., data = mimivc_icu_cohort_train) |>  
  # Mean imputation for normal variables  
  step_impute_mean(bicarbonate, chloride, hematocrit, sodium, heart_rate) |>  
  
  # Median imputation for skewed variables  
  step_impute_median(creatinine, glucose, potassium, wbc,  
    non_invasive_blood_pressure_diastolic,  
    non_invasive_blood_pressure_systolic,  
    respiratory_rate, temperature_fahrenheit) |>  
  
  # KNN imputation for categorical variables
```

```
step_impute_knn(marital_status) |>

# create dummy variable
step_dummy(all_nominal_predictors()) |>

# zero-variance filter
step_zv(all_nominal_predictors()) |>
print()
```

---

## — Recipe

### — Inputs

Number of variables by role

outcome: 1  
predictor: 18

### — Operations

- Mean imputation for: bicarbonate, chloride, hematocrit, sodium, ...
- Median imputation for: creatinine, glucose, potassium, wbc, ...
- K-nearest neighbor imputation for: marital\_status
- Dummy variables from: all\_nominal\_predictors()
- Zero variance filter on: all\_nominal\_predictors()

```
stacks_recipe
```

---

## — Recipe

### — Inputs

Number of variables by role

outcome: 1  
predictor: 18

### — Operations

- Mean imputation for: bicarbonate, chloride, hematocrit, sodium, ...

- Median imputation for: creatinine, glucose, potassium, wbc, ...
- K-nearest neighbor imputation for: marital\_status
- Dummy variables from: all\_nominal\_predictors()
- Zero variance filter on: all\_nominal\_predictors()

```
set.seed(203)
folds <- vfold_cv(mimiciv_icu_cohort_train, v = 2)
```

```
logit_mod <- logistic_reg(
  penalty = tune(),
  mixture = tune()
) |>
  set_engine("glmnet", standardize = TRUE)

logit_wf <- workflow() |>
  add_recipe(stacks_recipe) |>
  add_model(logit_mod)

logit_grid <- grid_regular(
  penalty(range = c(-4, 1)),
  mixture(),
  levels = c(10, 3)
)

logit_res <-
  tune_grid(
    object = logit_wf,
    resamples = folds,
    grid = logit_grid,
    control = control_stack_grid()
  )
```

**i** The workflow being saved contains a recipe, which is 5.84 Mb in **i** memory. If this was not intentional, please set the control setting **i** `save\_workflow = FALSE`.

```
logit_res
```

```
# Tuning results
# 2-fold cross-validation
# A tibble: 2 × 5
  splits          id    .metrics      .notes      .predictions
  <list>         <chr> <list>      <list>      <list>
1 <split [23610/23611]> Fold1 <tibble [90 × 6]> <tibble [0 × 3]> <tibble>
2 <split [23611/23610]> Fold2 <tibble [90 × 6]> <tibble [0 × 3]> <tibble>
```



```

rf_mod <- rand_forest(
  mode = "classification",
  mtry = tune(), # number of predictors randomly sampled in each split
  trees = tune() # number of trees in ensemble
) |>
  set_engine("ranger")

rf_wf <- workflow() |>
  add_recipe(stacks_recipe) |>
  add_model(rf_mod)

rf_grid <- grid_regular(
  trees(range = c(200L, 1000L)),
  mtry(range = c(1L, 8L)),
  levels = c(2, 2)
)

rf_res <- tune_grid(
  object = rf_wf,
  resamples = folds,
  grid = rf_grid,
  control = control_stack_grid()
)

```

**i** The workflow being saved contains a recipe, which is 5.84 Mb in **i** memory. If this was not intentional, please set the control setting **i** `save\_workflow = FALSE`.

```
rf_res
```

```

# Tuning results
# 2-fold cross-validation
# A tibble: 2 × 5
  splits          id    .metrics          .notes          .predictions
  <list>         <chr> <list>          <list>          <list>
1 <split [23610/23611]> Fold1 <tibble [12 × 6]> <tibble [0 × 3]> <tibble>
2 <split [23611/23610]> Fold2 <tibble [12 × 6]> <tibble [0 × 3]> <tibble>

```

```

gb_mod <- boost_tree(
  mode = "classification",
  trees = 1000,
  tree_depth = tune(),
  learn_rate = tune()
) |>
  set_engine("xgboost")
gb_mod

```

Boosted Tree Model Specification (classification)

Main Arguments:

```
trees = 1000  
tree_depth = tune()  
learn_rate = tune()
```

Computational engine: xgboost

```
gb_wf <- workflow() |>  
  add_recipe(stacks_recipe) |>  
  add_model.gb_mod()  
  
gb_grid <- grid_regular(  
  tree_depth(range = c(3L, 10L)),  
  learn_rate(range = c(0.01, 0.3)),  
  levels = c(1, 2)  
)  
  
gb_res <- tune_grid(  
  object = gb_wf,  
  resamples = folds,  
  grid = gb_grid,  
  control = control_stack_grid()  
)
```

**i** The workflow being saved contains a recipe, which is 5.84 Mb in **i** memory. If this was not intentional, please set the control setting **i** `save\_workflow = FALSE`.

```
gb_res
```

# Tuning results

# 2-fold cross-validation

# A tibble: 2 × 5

	splits	id	.metrics	.notes	.predictions
	<list>	<chr>	<list>	<list>	<list>
1	<split [23610/23611]>	Fold1	<tibble [6 × 6]>	<tibble [0 × 3]>	<tibble>
2	<split [23611/23610]>	Fold2	<tibble [6 × 6]>	<tibble [0 × 3]>	<tibble>

```
model_st <- stacks() |>  
  add_candidates(logit_res) |>  
  add_candidates(rf_res) |>  
  add_candidates(gb_res) |>  
  # determine how to combine their predictions  
  blend_predictions(  
    penalty = 10^(-6:2),  
    metrics = c("roc_auc")  
  ) |>  
  # fit candidates with nonzero stacking coefficients  
  fit_members()
```

Warning: Predictions from 26 candidates were identical to those from existing candidates and were removed from the data stack.

Warning: The `...` are not used in this function but one or more arguments were passed: 'metrics'

```
model_st
```

— A stacked ensemble model —————

Out of 23 possible candidate members, the ensemble retained 7.

Penalty: 0.001.

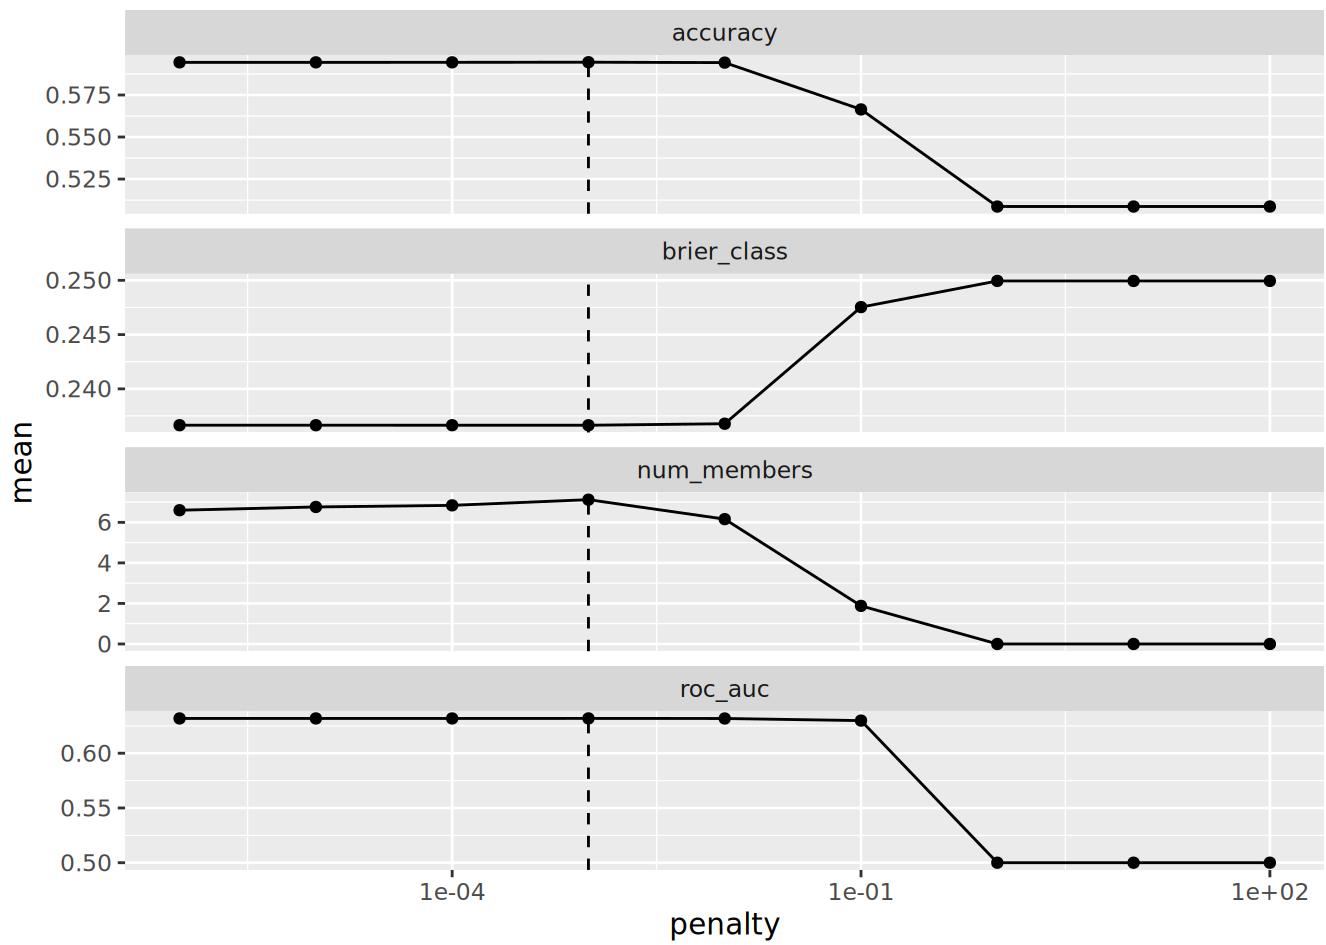
Mixture: 1.

The 7 highest weighted member classes are:

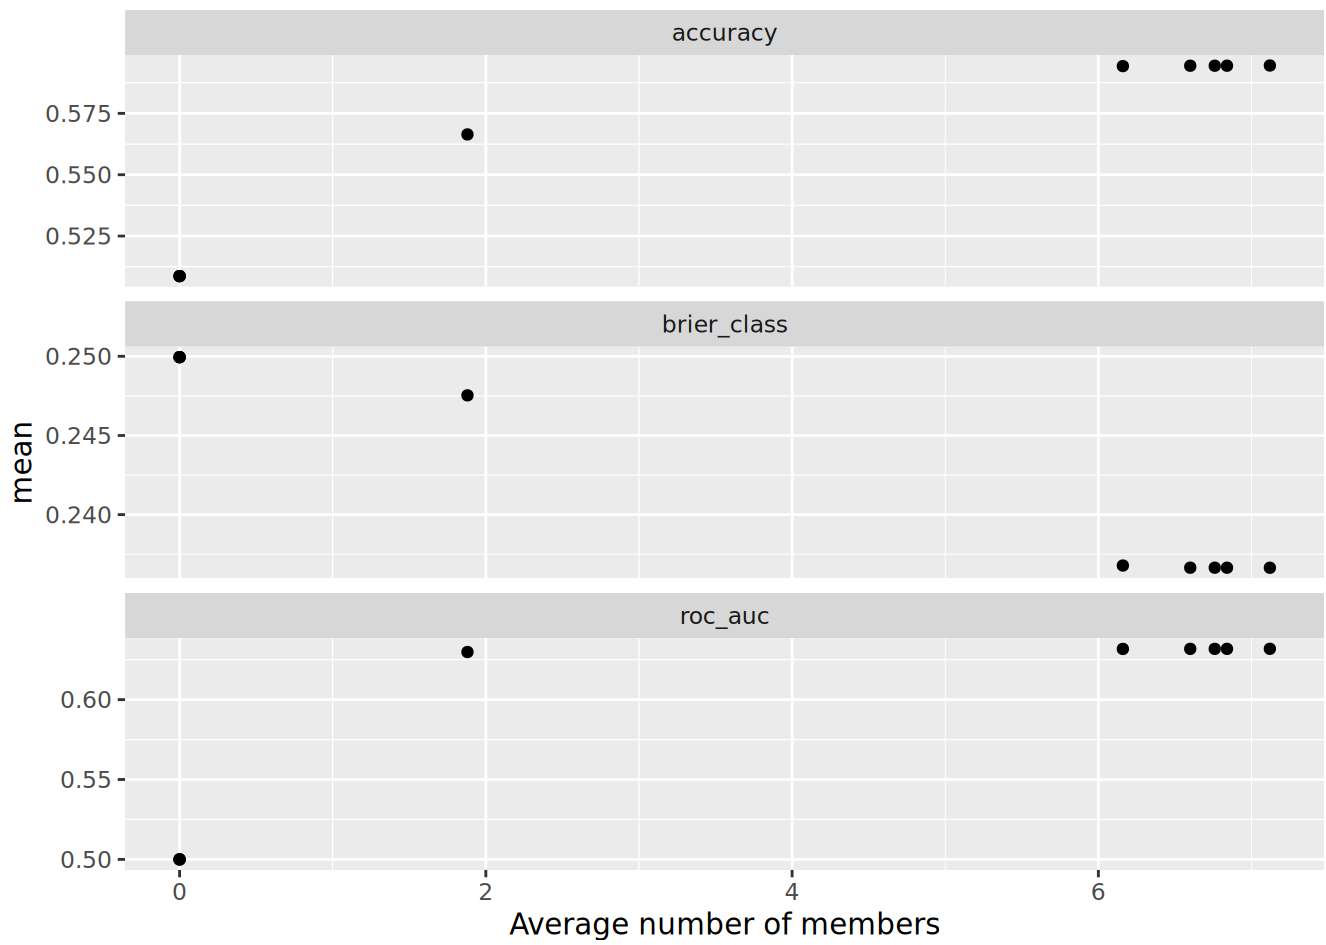
```
# A tibble: 7 × 3
```

	member	type	weight
	<chr>	<chr>	<dbl>
1	.pred_FALSE_rf_res_1_2	rand_forest	5.03
2	.pred_FALSE_rf_res_1_4	rand_forest	2.15
3	.pred_FALSE_logit_res_1_21	logistic_reg	0.376
4	.pred_FALSE_rf_res_1_1	rand_forest	0.326
5	.pred_FALSE_rf_res_1_3	rand_forest	0.317
6	.pred_FALSE_gb_res_1_1	boost_tree	0.119
7	.pred_FALSE_logit_res_1_11	logistic_reg	0.00994

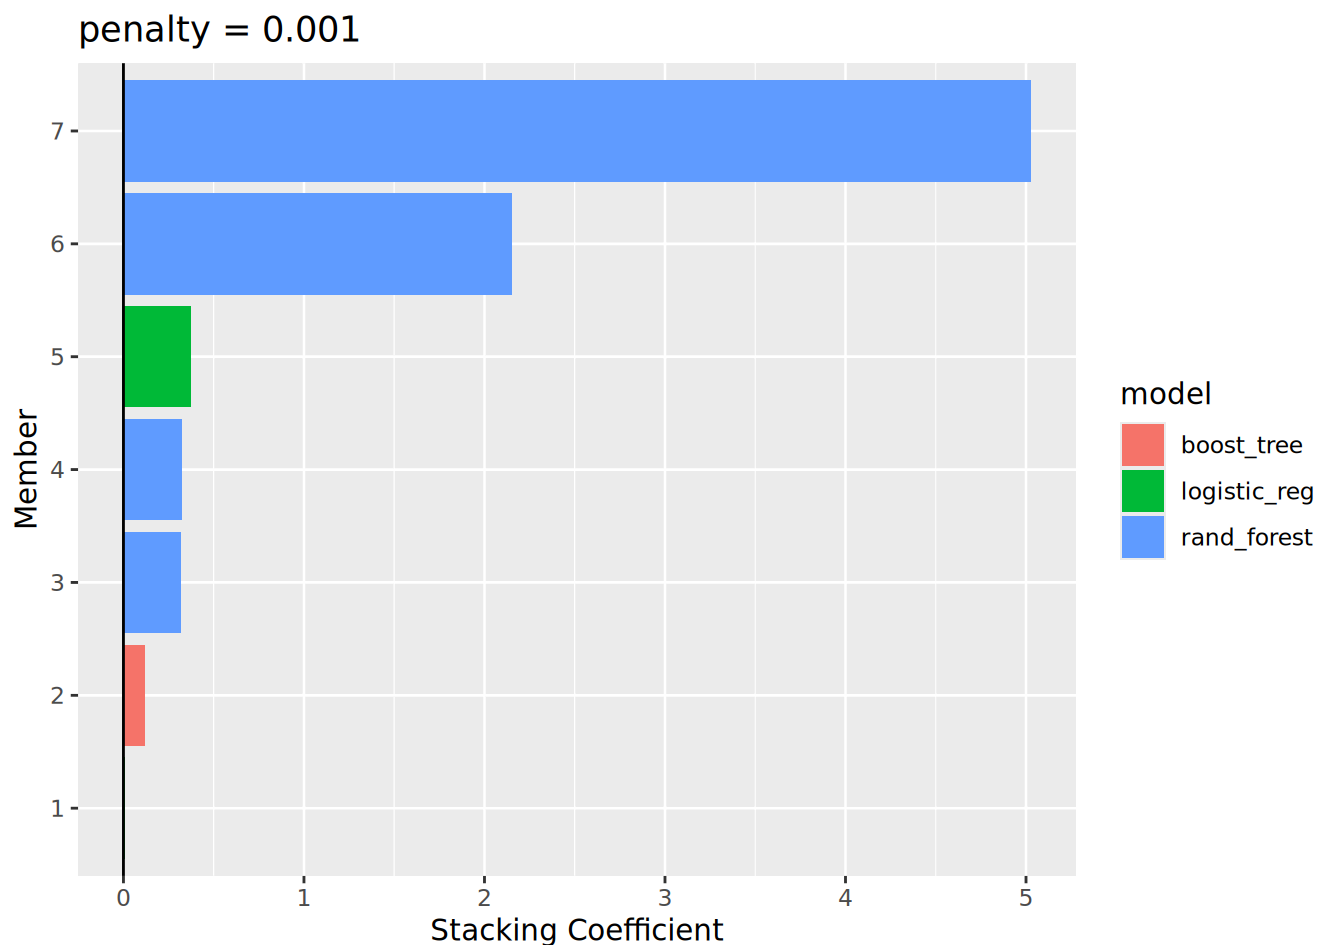
```
autoplot(model_st)
```



```
autoplot(model_st, type = "members")
```



```
autoplot(model_st, type = "weights")
```



```
collect_parameters(model_st, "rf_res")
```

# A tibble: 4 × 5

	member	mtry	trees	terms	coef
	<chr>	<int>	<int>	<chr>	<dbl>
1	rf_res_1_1	1	200	.pred_FALSE_rf_res_1_1	0.326
2	rf_res_1_2	1	1000	.pred_FALSE_rf_res_1_2	5.03
3	rf_res_1_3	8	200	.pred_FALSE_rf_res_1_3	0.317
4	rf_res_1_4	8	1000	.pred_FALSE_rf_res_1_4	2.15

```
mimic_pred <- mimivicv_icu_cohort_test %>%
  bind_cols(predict(model_st, ., type = "prob")) %>%
  print(width = Inf)
```

# A tibble: 47,223 × 21

	los_long	gender	age_intime	marital_status	race
	<fct>	<fct>	<int>	<fct>	<fct>
1	FALSE	F	52	WIDOWED	white
2	FALSE	F	46	MARRIED	white
3	FALSE	F	57	SINGLE	other
4	TRUE	M	56	<NA>	other
5	FALSE	F	83	MARRIED	white

6	TRUE	F	82	MARRIED	white		
7	TRUE	F	81	WIDOWED	white		
8	TRUE	M	90	WIDOWED	white		
9	TRUE	M	53	SINGLE	white		
10	FALSE	F	58	<NA>	white		
	first_careunit				bicarbonate	chloride	
	<fct>				<dbl>	<dbl>	
1	Medical Intensive Care Unit (MICU)				25	95	
2	Medical/Surgical Intensive Care Unit (MICU/SICU)				NA	98	
3	Cardiac Vascular Intensive Care Unit (CVICU)				24	102	
4	Other				18	NA	
5	Medical Intensive Care Unit (MICU)				26	85	
6	Medical/Surgical Intensive Care Unit (MICU/SICU)				23	98	
7	Medical Intensive Care Unit (MICU)				27	111	
8	Other				23	102	
9	Other				18	106	
10	Cardiac Vascular Intensive Care Unit (CVICU)				NA	NA	
	creatinine	glucose	hematocrit	potassium	sodium	wbc	heart_rate
	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>
1	0.7	102	41.1	6.7	126	6.9	91
2	NA	NA	NA	4.1	139	NA	86
3	0.9	288	34.9	3.5	137	7.2	80
4	3.1	95	34.3	6.5	125	16.8	110.
5	1.4	133	22.4	5.7	120	9.8	114
6	2.8	117	25.5	4.9	135	17.9	91
7	0.6	173	34.7	4.4	144	10.5	106.
8	1.9	105	29.9	4.4	140	5.1	93.5
9	0.9	269	43.1	5.3	135	16.9	106
10	NA	NA	NA	NA	NA	NA	80
	non_invasive_blood_pressure_diastolic				non_invasive_blood_pressure_systolic		
				<dbl>			<dbl>
1				48			84
2				56			73
3				62			98.5
4				80			112
5				65			109
6				51			118
7				51			102
8				61			108
9				99			140
10				72			109
	respiratory_rate	temperature_fahrenheit	.pred_TRUE	.pred_FALSE			
	<dbl>	<dbl>	<dbl>	<dbl>			
1	24	98.7	0.575	0.425			
2	19	97.7	0.470	0.530			
3	14	97.2	0.436	0.564			
4	21	97.9	0.643	0.357			
5	24	97.7	0.571	0.429			
6	18	96.9	0.547	0.453			
7	25	98.6	0.522	0.478			
8	22.5	98.1	0.523	0.477			

9	12	96.7	0.487	0.513
10	17	99	0.611	0.389

# i 47,213 more rows

## 4. Compare model classification performance on the test set. [🔗](#)

Report both the area under ROC curve and accuracy for each machine learning algorithm and the model stacking. Interpret the results. What are the most important features in predicting long ICU stays? How do the models compare in terms of performance and interpretability?

### Logistic Regression

```
logit_fit_final |>
  collect_metrics()
```

```
# A tibble: 3 × 4
  .metric      .estimator .estimate .config
  <chr>        <chr>         <dbl> <chr>
1 accuracy    binary         0.567 Preprocessor1_Model1
2 roc_auc     binary         0.592 Preprocessor1_Model1
3 brier_class binary         0.244 Preprocessor1_Model1
```

The best logistic regression model has accuracy of 0.567 and roc\_auc of 0.59. The results shows that 56.7% of ICU stay length are correctly classified by the model. The ROC AUC measures the model's ability to distinguish between classes. The score of 0.59 suggests that the model performs slightly better than random guessing (0.5).

### Random Forest

```
rf_fit_final |>
  collect_metrics()
```

```
# A tibble: 3 × 4
  .metric      .estimator .estimate .config
  <chr>        <chr>         <dbl> <chr>
1 accuracy    binary         0.590 Preprocessor1_Model1
2 roc_auc     binary         0.629 Preprocessor1_Model1
3 brier_class binary         0.237 Preprocessor1_Model1
```

The accuracy of random forest model is 0.58, meaning 58% of models are correctly predicted. The ROC AUC is 0.62, which means 62% of the time the model distinguish the classes successfully.

### XGBoost

```
gb_fit_final |>
  collect_metrics()
```

```
# A tibble: 3 × 4
  .metric      .estimator .estimate .config
```



	<chr>	<chr>	<dbl>	<chr>
1	accuracy	binary	0.560	Preprocessor1_Model1
2	roc_auc	binary	0.580	Preprocessor1_Model1
3	brier_class	binary	0.282	Preprocessor1_Model1

The accuracy of XGBoost model is 0.56%, meaning 56% of los\_long are correctly predicted. The ROC AUC of 0.58 means that 58% of the time the model distinguish the classes successfully.

## Model Stacking

```
yardstick::roc_auc(
  mimic_pred,
  truth = los_long,
  contains(".pred_FALSE")
)
```

```
# A tibble: 1 × 3
  .metric .estimator .estimate
  <chr>    <chr>        <dbl>
1 roc_auc binary      0.370
```

The ROC\_AUC of model stacking is 0.3687729. which shows that it's not successfully predicted the los\_long.

## Most important features

```
logit_fit_final |>
  extract_fit_parsnip() |>
  tidy() |>
  arrange(desc(estimate))
```

```
# A tibble: 27 × 3
  term                                estimate penalty
  <chr>                                <dbl>    <dbl>
1 first_careunit_Medical.Surgical.Intensive.Care.Unit..MICU.S... 0.160    0.00168
2 hematocrit                        0.0977   0.00168
3 non_invasive_blood_pressure_systolic 0.0974   0.00168
4 first_careunit_Medical.Intensive.Care.Unit..MICU.             0.0942   0.00168
5 chloride                          0.0915   0.00168
6 (Intercept)                     0.0374   0.00168
7 marital_status_WIDOWED          0.0172   0.00168
8 race_black                       0.000442 0.00168
9 bicarbonate                       0         0.00168
10 glucose                          0         0.00168
# i 17 more rows
```

Based on the best logistic model, the most important features with the largest estimates are first\_careunit\_Medical.Surgical.Intensive.Care.Unit..MICU.SICU., hematocrit, non\_invasive\_blood\_pressure\_systolic, first\_careunit\_Medical.Intensive.Care.Unit..MICU., and chloride.

## Compare performance and interpretability

Comparing the accuracy and roc auc of 4 models, the first 3 models, logistic regression, random forest, and XGBoost, have similar performance in accuracy. Random Forest has the highest roc auc among the 3 models. The stacking model performed poorly, which an ROC AUC below 0.5, indicating it performs worse than random guessing. This suggests issues with the model blending or that the individual models may not complement each other well.

**Logistic regression** is the most interpretable among all 4 models. The coefficients directly show the effect of each feature on the outcome. It's excellent for understanding relationships in data. **Random Forest** provides variable importance, but understanding individual predictions can be difficult. **XGBoost** is even more complex, though feature importance and SHAP values can offer insights into how predictions are made. **Model Stacking** is the hardest to interpret since it's a combination of other models, making it a "black-box" approach.